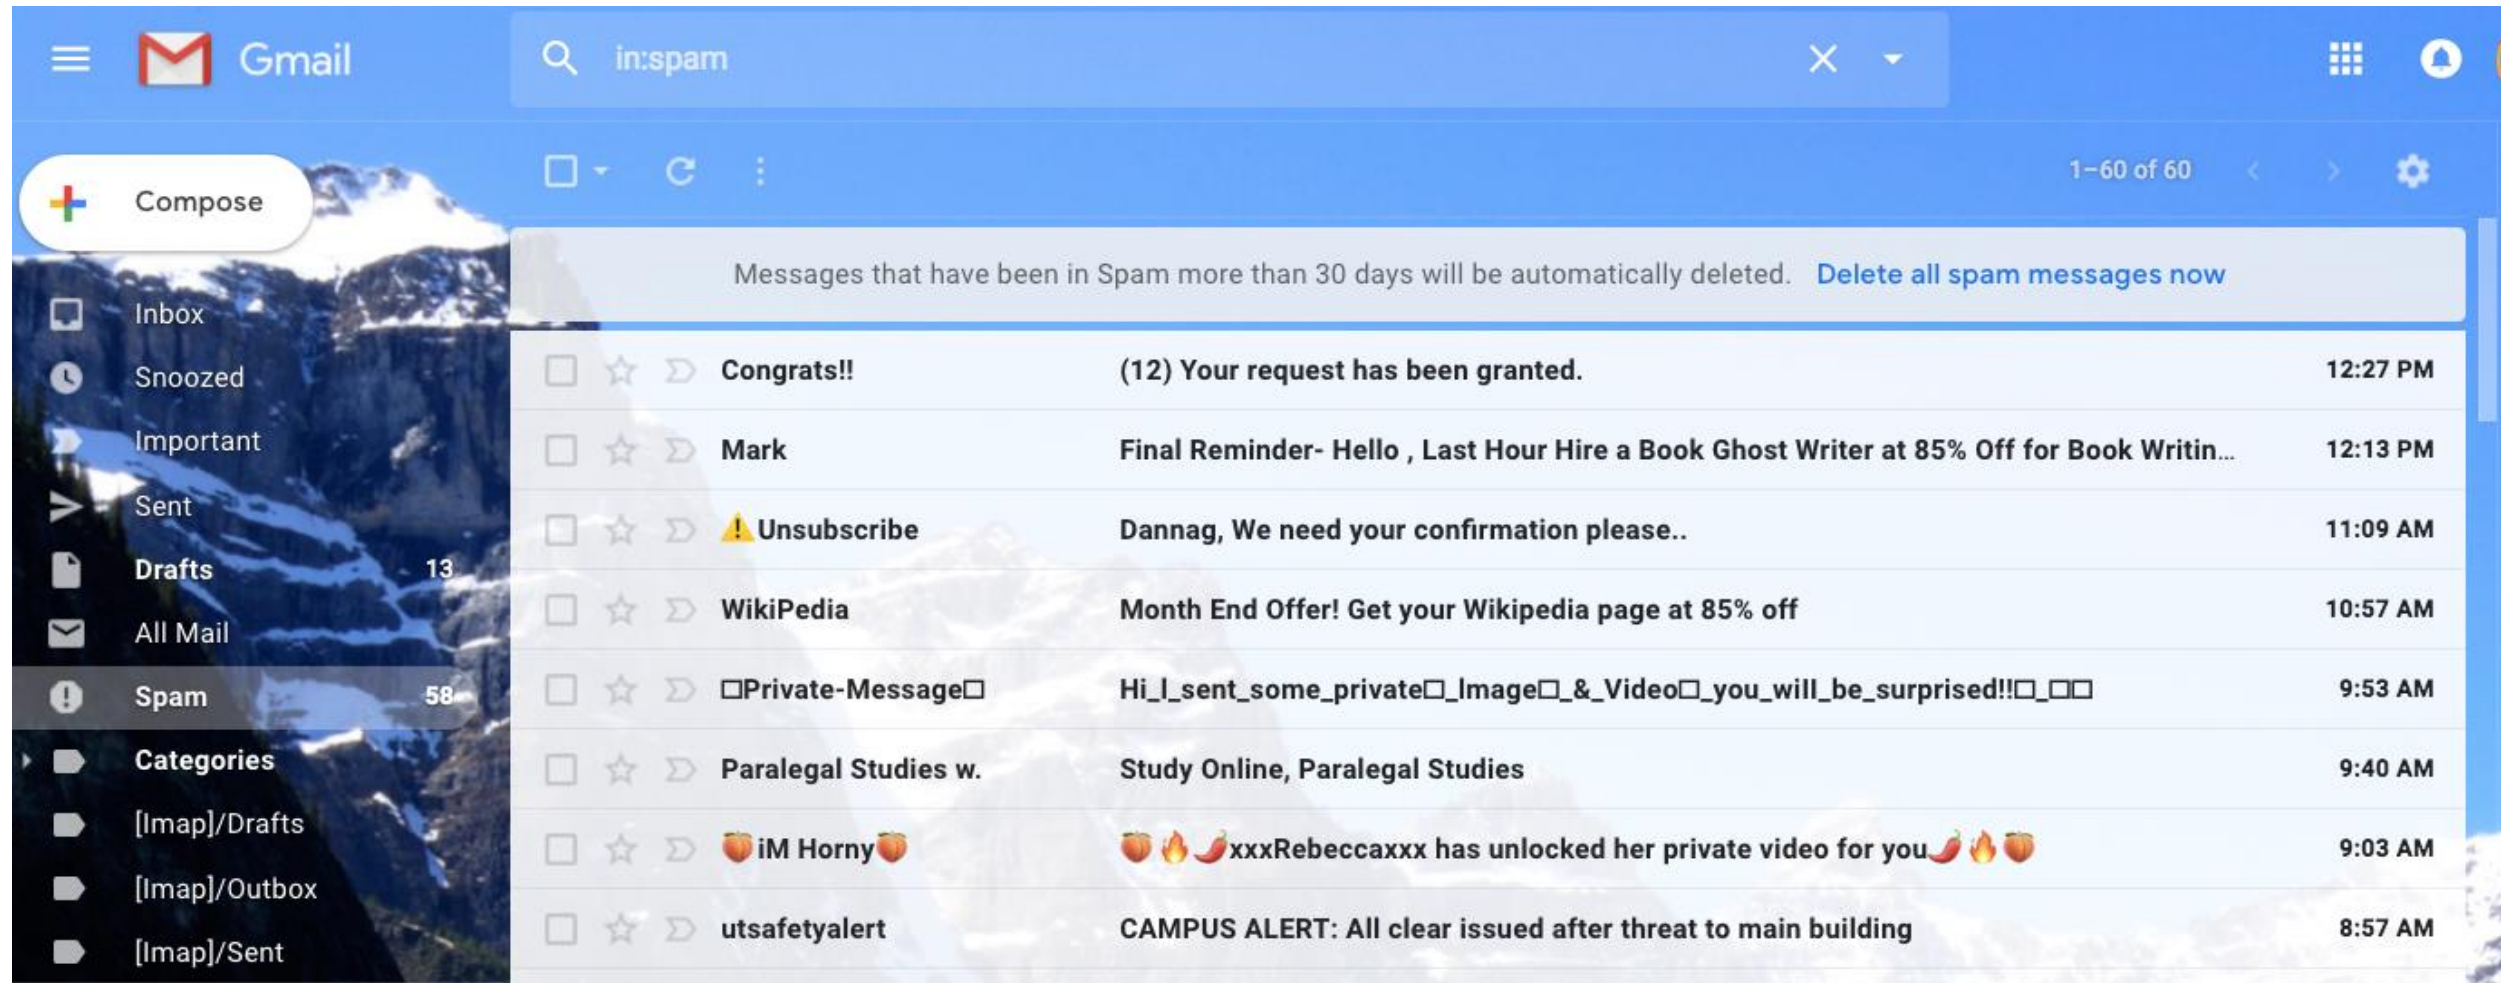


Artificial Neural Networks


Binary Classification

Distinguish 2 classes

Binary Classification: Spam Detection



Binary Classification: Resume Pre-Screening




Whaii

**AI AND BLOCKCHAIN
STAFFING & RECRUITMENT**

- Decentralizing and simplifying the staffing industry
- Eliminate the intermediaries between job seeker, through an open ecosystem of hiring managers by using Blockchain, AI, and other technologies, that ultimately makes hiring more cost-effective

[BUY COIN](#) [WHITE PAPER](#)



Prime Talent Chain
Hiring Decentralized


[ABOUT](#) [PLATFORM](#) [ROADMAP](#) [COIN](#) [TEAM](#) [WHITEPAPER](#) [SIGN IN](#) [BUY COIN](#)

Automatically screen resumes


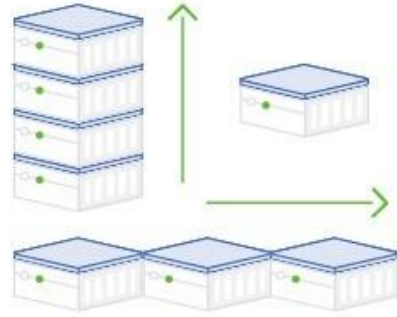
Trained with over 20 million diverse profiles, Skillate's AI algorithm helps to screen and shortlist resume with the click of a button. Seamlessly integrate with all external channels and ATS to source resume directly

Matching algorithm and candidate recommendation

Skillate's matching engine maps all the relevant profiles with the job requirements - be it skills, education or experience and recommends the best candidate



HOME [PRODUCT](#) [ABOUT US](#) [BLOG](#) [REQUEST A DEMO](#)

Binary Classification: Cancer Diagnosis



The image shows the top portion of the PathAI website. At the top is a navigation bar with the PathAI logo on the left and links for 'What we do', 'About us', 'News', 'Careers', 'Pathologists', and 'Partner with Us' in the center. A red 'Partner Login' button is on the right. Below the navigation bar is a large circular image of a histological slide with green and yellow highlights. To the right of the image is the heading 'Pathology Evolved.' followed by the text 'Advanced learning toward faster, more accurate diagnosis of disease.'

PathAI

What we do About us News Careers Pathologists Partner with Us

Partner Login

Pathology Evolved.

Advanced learning toward faster,
more accurate diagnosis of disease.


Binary Classification: Cognitive Impairment Recognition by Apple App Usage



Image Credit: <https://www.techradar.com/news/the-10-best-phones-for-seniors>

<https://www.technologyreview.com/f/615032/the-apps-you-use-on-your-phone-could-help-diagnose-your-cognitive-health/>


Binary Classification: Sentiment Analysis



[What's the Tomatometer®?](#)
[Critics](#)
[SIGN UP](#)
[LOG IN](#)


[MOVIES](#)
[TV SHOWS](#)
[RT PODCAST](#)
[NEWS](#)
[SHOWTIMES](#)

CRITIC REVIEWS FOR *MULAN*




Its cast, its attitude, its overall eagerness to please – all benefits, one would think – don't add up to a good movie. They add up to a blueprint of the movie this ought to be.

October 23, 2020 | Rating: 2.5/5 | [Full Review...](#)




K. Austin Collins
Rolling Stone
★ TOP CRITIC



While glorious to look at, the movie still feels slightly hollow. All the right pieces are there, but an emotional connection to the characters is lacking.

September 10, 2020 | Rating: 6.8/10 | [Full Review...](#)



Amy Amatangelo
Paste Magazine
★ TOP CRITIC

Binary Classification: Food Quality Control



Machine Learning: Using Algorithms to Sort Fruit

Demo: <https://www.youtube.com/watch?v=Bl3XzBWpZbY>

Goal: Design Models that Generalize Well to New, Previously Unseen Examples

Example:



Label:

Hairy

Hairy

Not Hairy



Hairy



Goal: Design Models that Generalize Well to New, Previously Unseen Examples

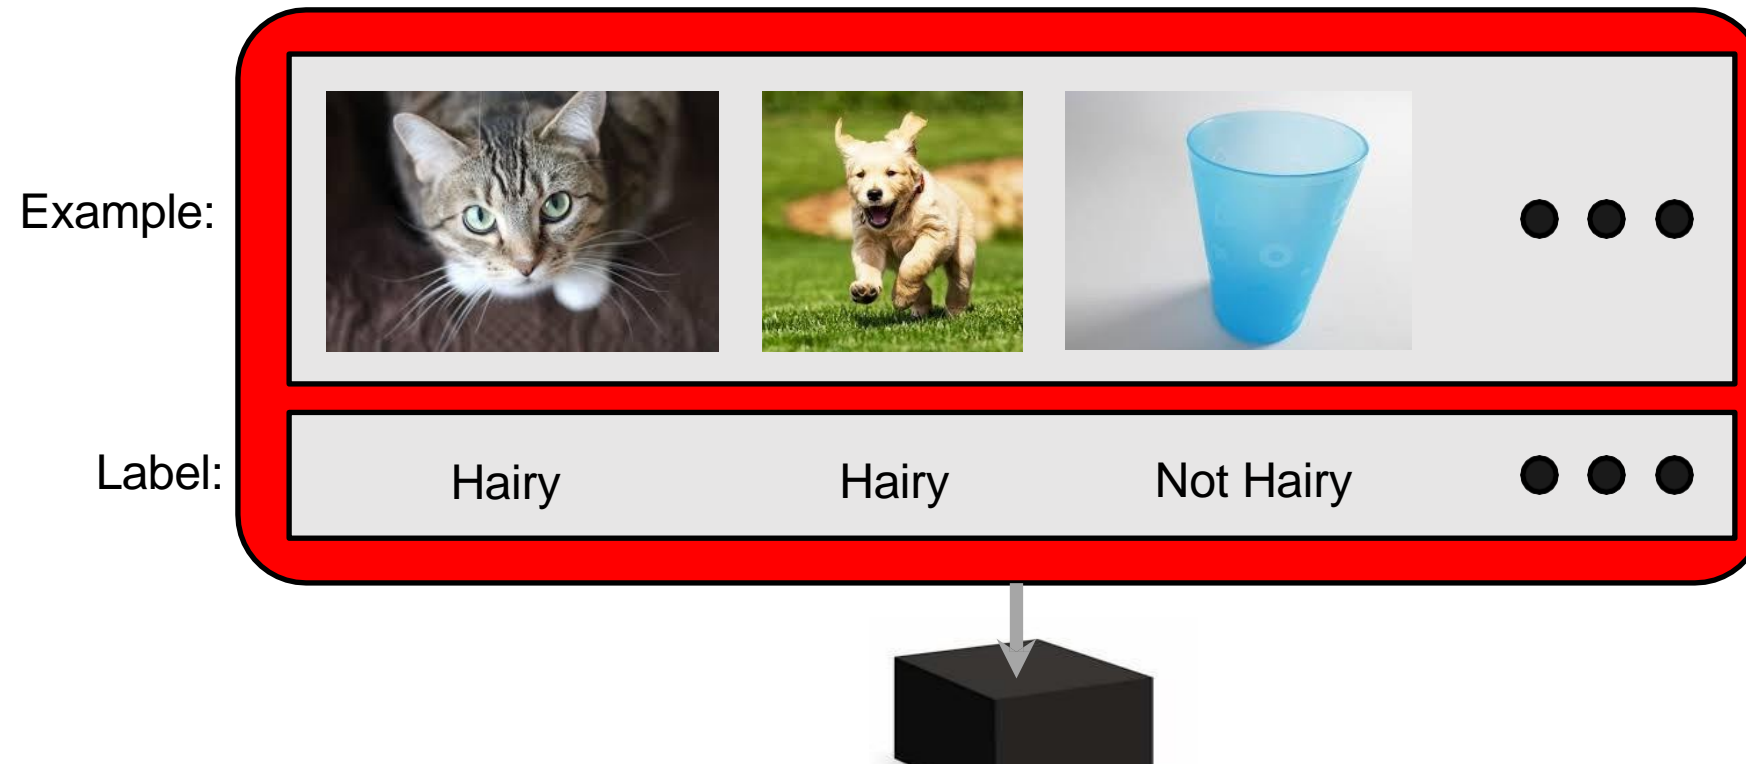
1. Split data into a “**training set**” and “**test set**”



Goal: Design Models that Generalize Well to New, Previously Unseen Examples


2. Train model on “**training set**” to try to minimize prediction error on it

Training Data




Goal: Design Models that Generalize Well to New, Previously Unseen Examples

3. Apply trained model on “**test set**” to measure generalization error





Example:




Goal: Design Models that Generalize Well to New, Previously Unseen Examples

3. Apply trained model on “**test set**” to measure generalization error




Example:




Goal: Design Models that Generalize Well to New, Previously Unseen Examples

3. Apply trained model on “**test set**” to measure generalization error



Example:

Test Data

	...
Label:	Hairy ...
Predicted Label:	Not Hairy ...

Inspiration: Animal's Computing Machinery

Neuron

- basic unit in the nervous system for receiving, processing, and transmitting information; e.g., messages such as...


“hot”




“loud”




“spicy”



Inspiration: Animal's Computing Machinery




Inspiration: Neuron “Firing”




- When the input signals exceed a certain threshold within a short period of time, a neuron “fires”
- Neuron “firing” (outputs signal) is an “all-or-none” process

Inspiration: Animal's Computing Machinery



<https://en.wikipedia.org/wiki/Nematode#/media/File:CelegansGoldsteinLabUNC.jpg>

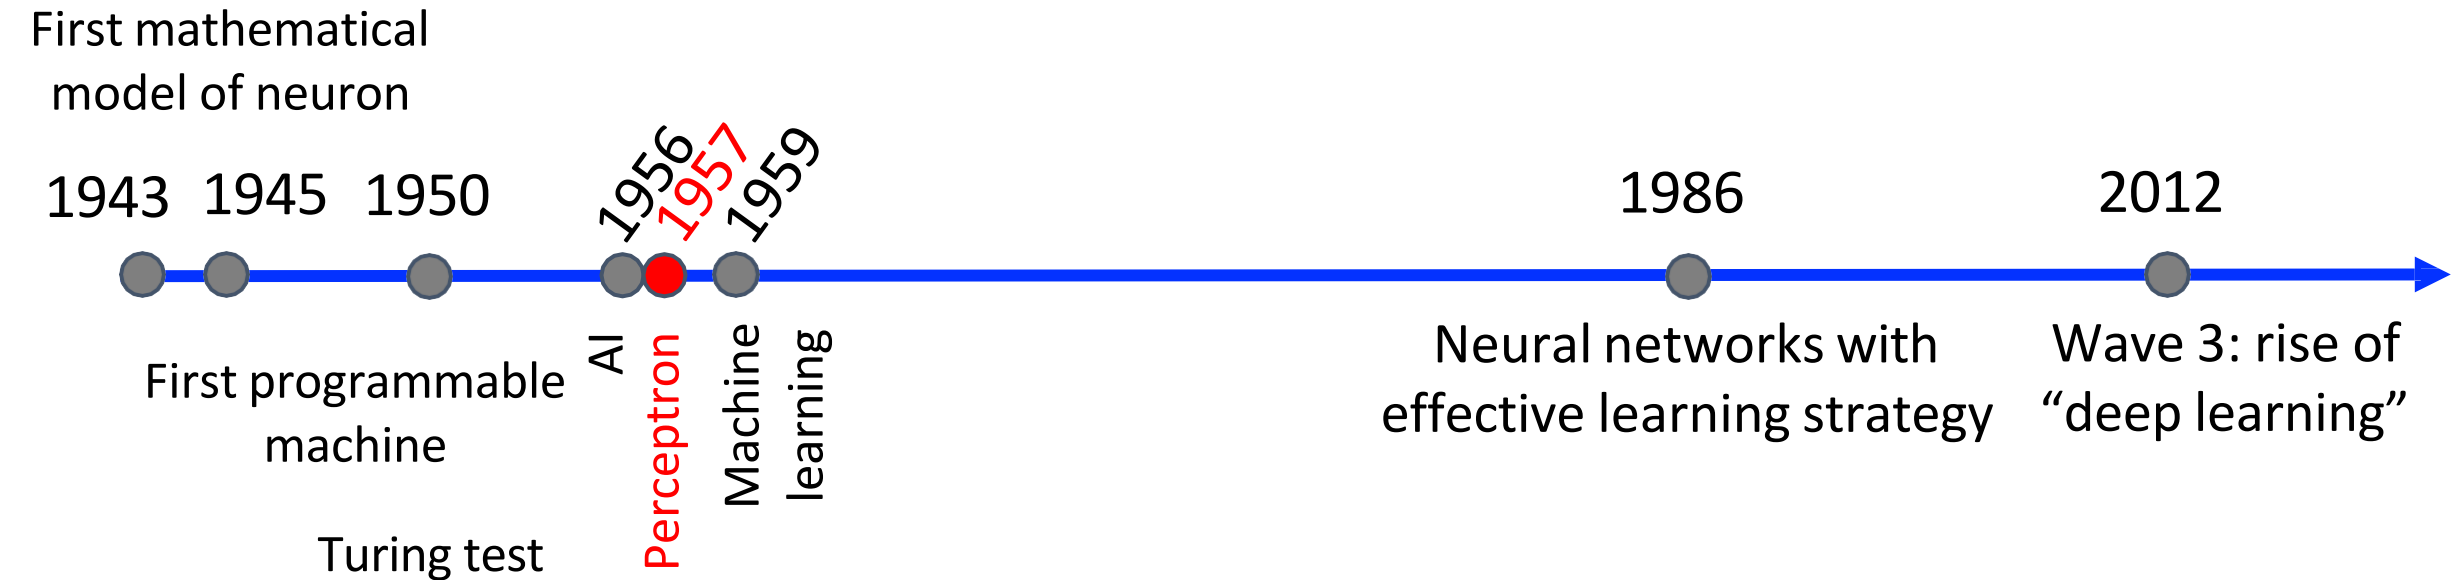
Nematode worm: 302 neurons



<https://www.britannica.com/science/human-nervous-system>

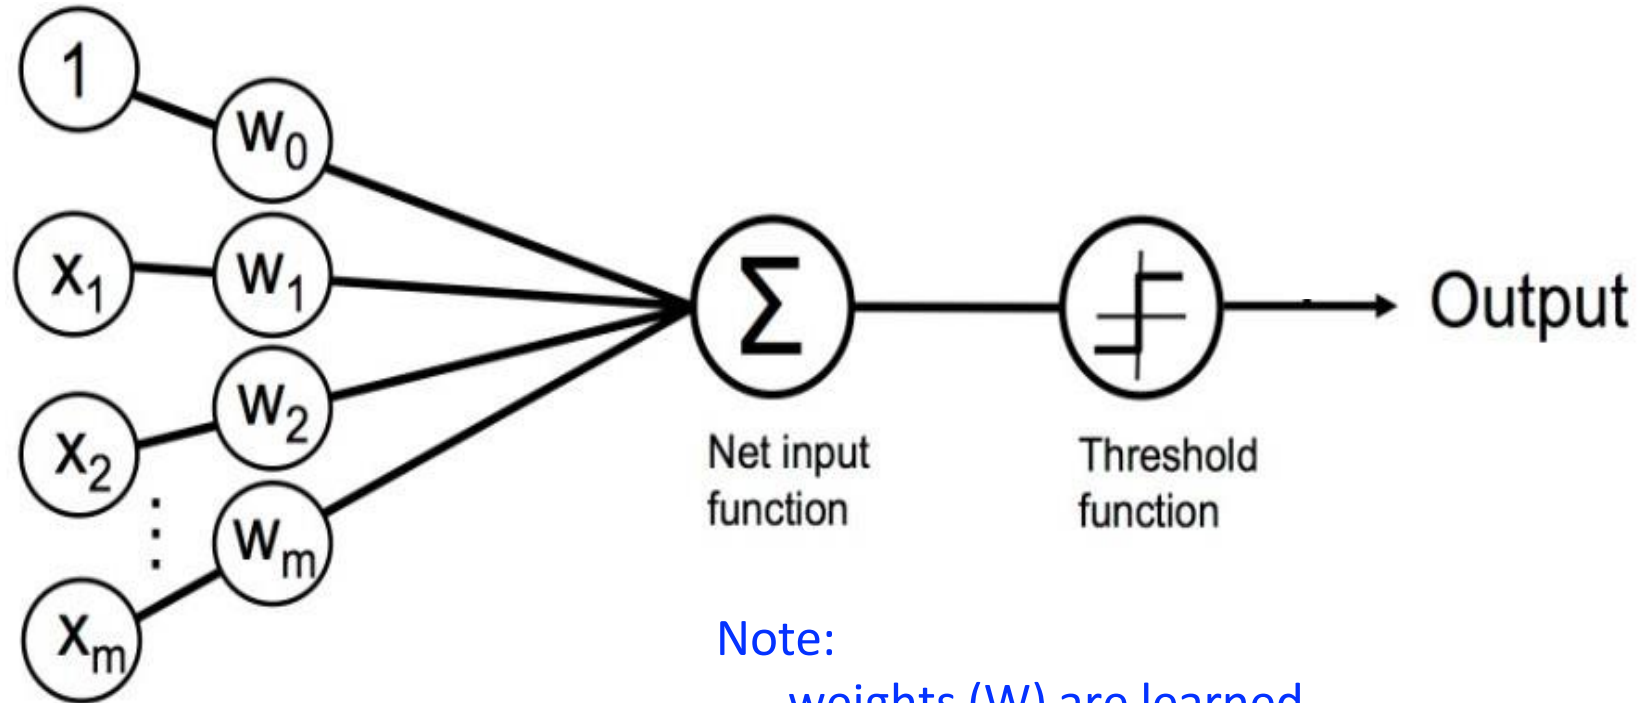
Human: ~Billions neurons

Historical Context: Artificial Neurons



Modern deep learning algorithms rely on techniques developed over the past 65 years.

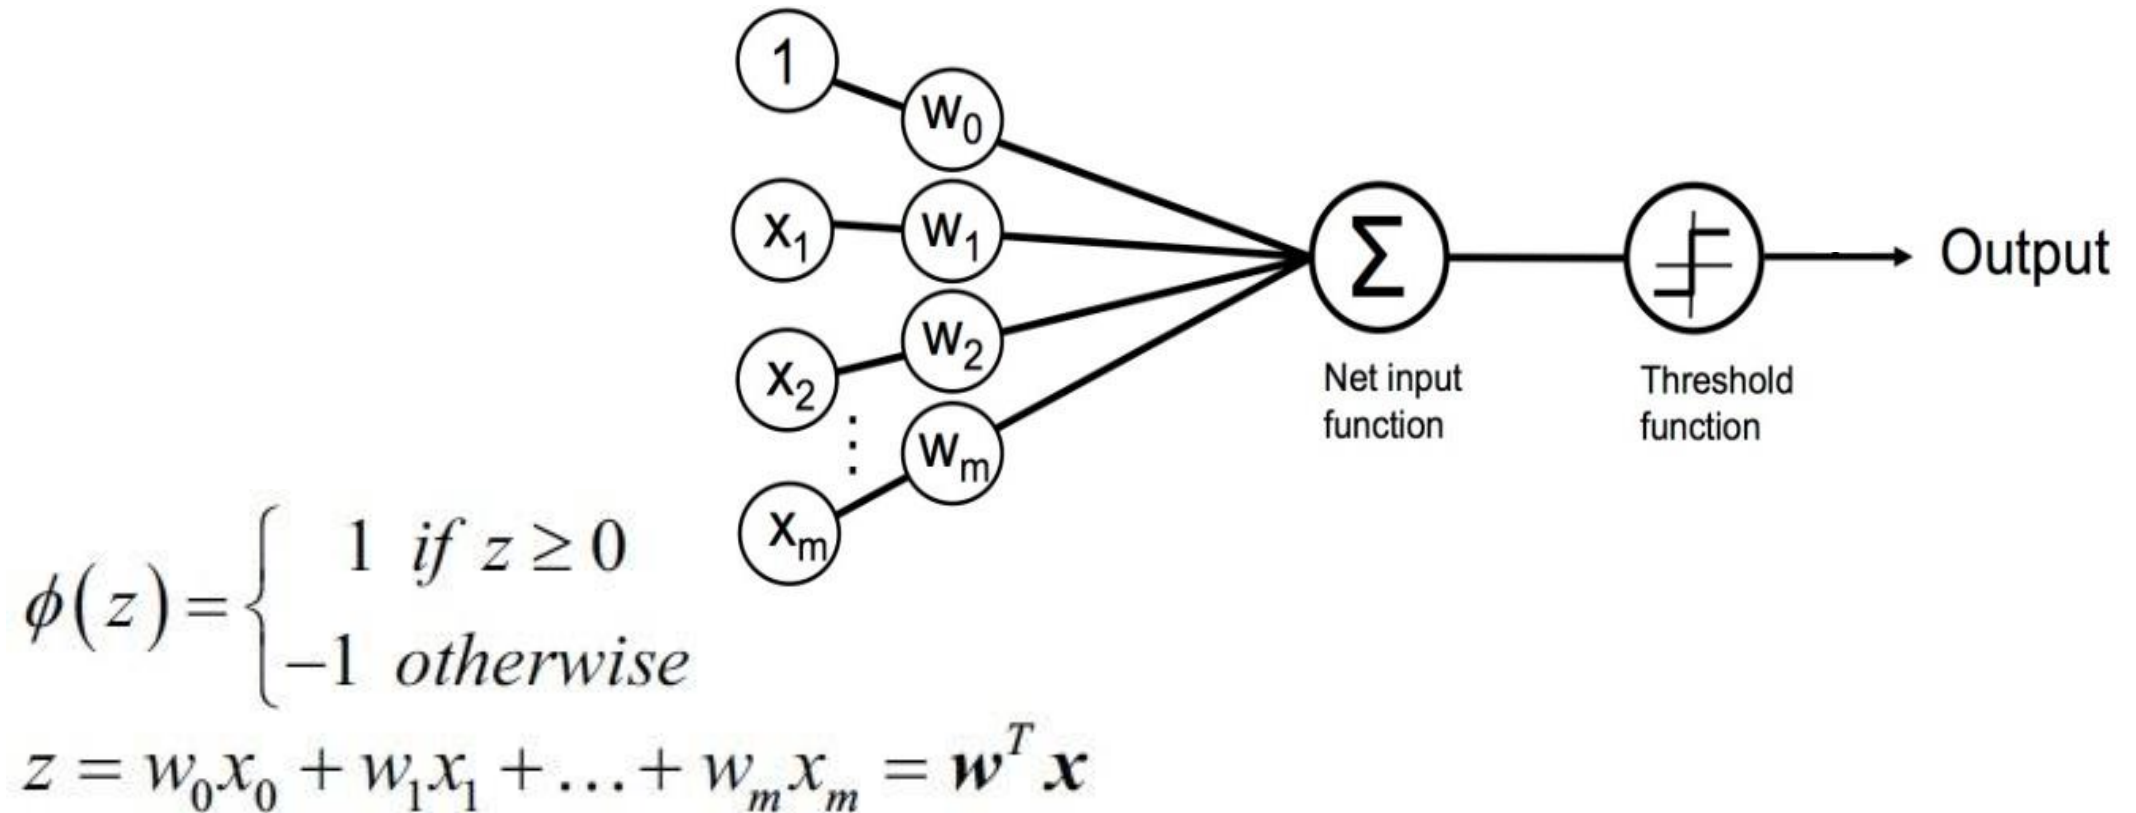
Perceptron: Model (Linear Threshold Unit)



Note:

- weights (W) are learned
- inputs and weights can be any value
- fires when combined input exceeds threshold

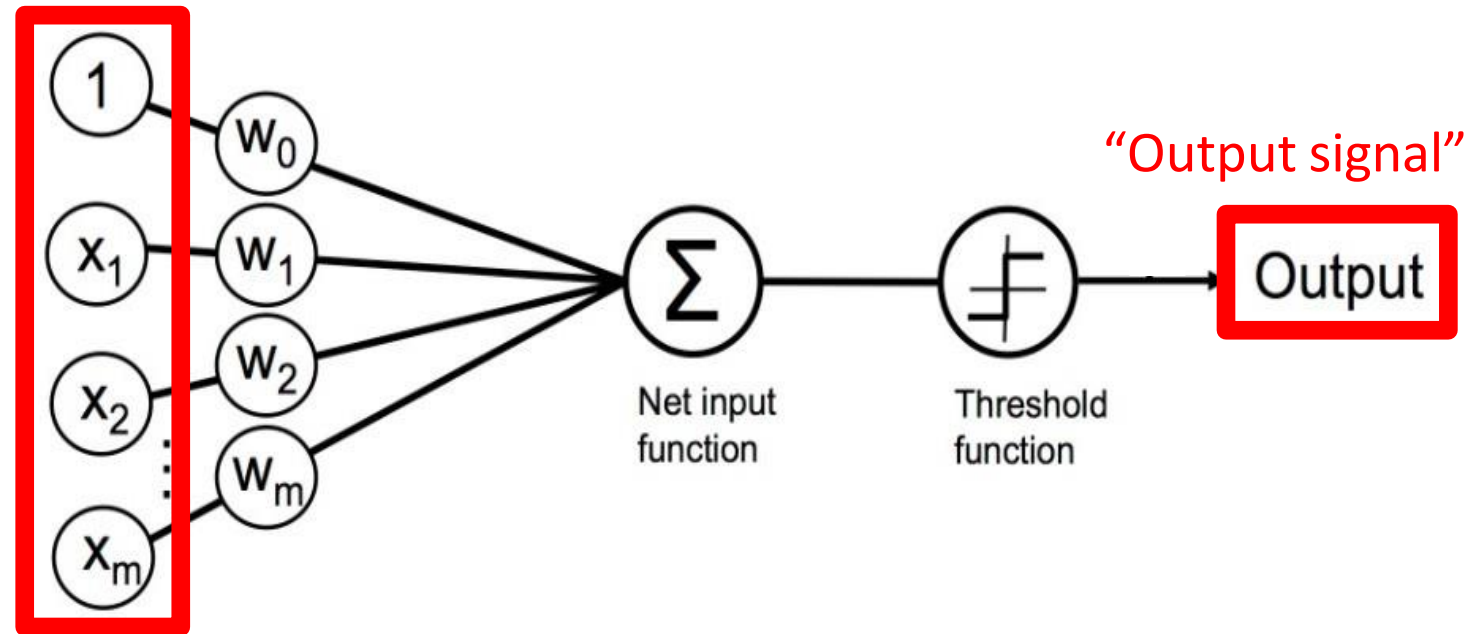
Perceptron: Model (Linear Threshold Unit)




Perceptron: Model (Linear Threshold Unit)

Artificial Neuron:

“Input signals”




Biological Neuron:



Python Machine Learning; Raschka & Mirjalili

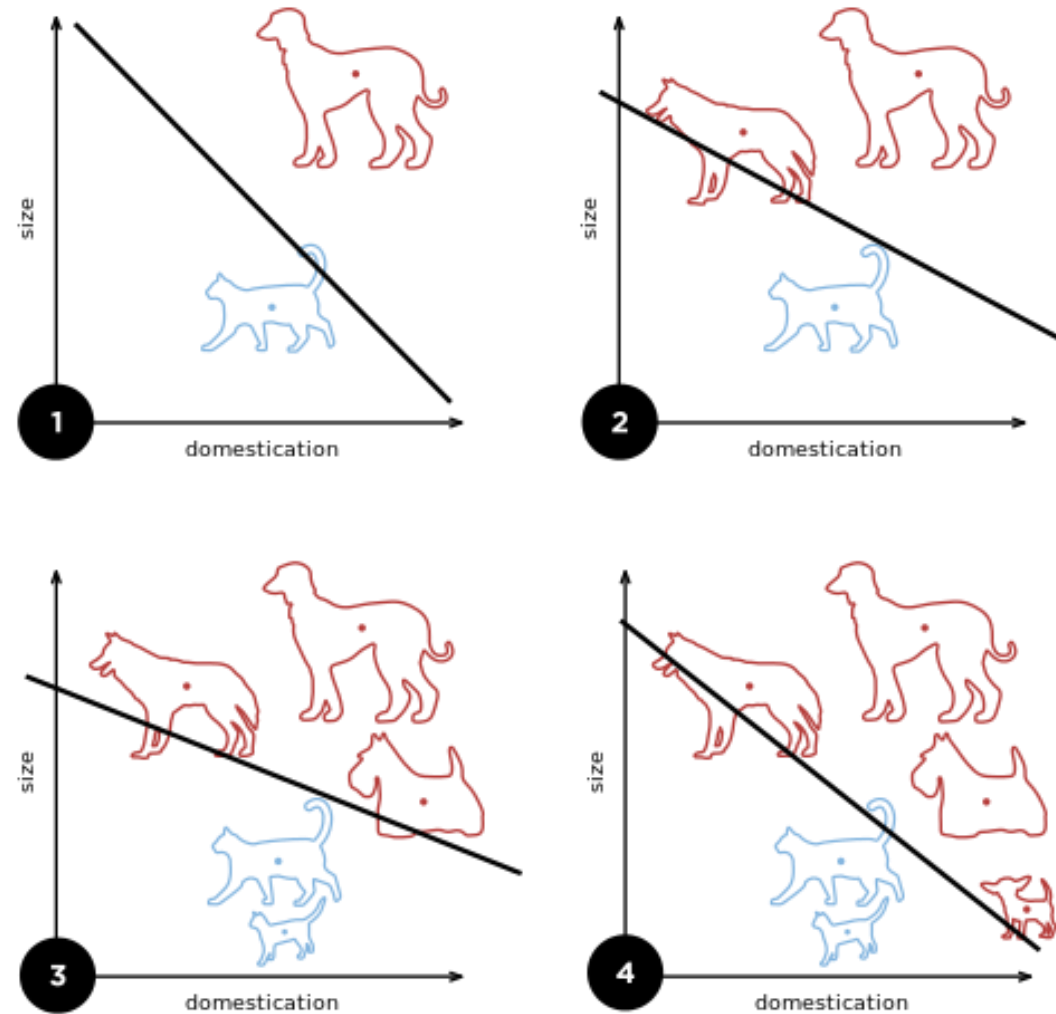
Perceptron: Learning Algorithm

Iteratively update linear boundary with observation of each additional example:



Perceptron: Learning Algorithm

Iteratively update linear boundary with observation of each additional example:



Perceptron: Learning Algorithm

1. Initialize weights to 0 or small random numbers
2. For each training sample:

1. Compute output value: $\sum_{j=0}^m \mathbf{x}_j \mathbf{w}_j = \mathbf{w}^T \mathbf{x}$

2. Update weights with the following definition: $\mathbf{w}_j := \mathbf{w}_j + \Delta \mathbf{w}_j$

$$\Delta \mathbf{w}_j = \eta (\text{target}^{(i)} - \text{output}^{(i)}) \mathbf{x}_j^{(i)}$$

Learning Rate

True Class Label

Predicted Class Label

Perceptron: Learning Algorithm - What Happens to Weights When It Predicts Correct Class Label?

1. Initialize weights to 0 or small random numbers
2. For each training sample:

1. Compute output value: $\sum_{j=0}^m \mathbf{x}_j \mathbf{w}_j = \mathbf{w}^T \mathbf{x}$

2. Update weights with the following definition: $\mathbf{w}_j := \mathbf{w}_j + \Delta \mathbf{w}_j$

equals 0, so no weight update

$$\Delta \mathbf{w}_j = \eta \left(\text{target}^{(i)} - \text{output}^{(i)} \right) \mathbf{x}_j^{(i)}$$

Learning Rate

↑

True Class Label

↑

Predicted Class Label

↑

Perceptron: Learning Algorithm - What Happens to Weights When It Predicts Wrong Class Label?

1. Initialize weights to 0 or small random numbers
2. For each training sample:

1. Compute output value: $\sum_{j=0}^m \mathbf{x}_j \mathbf{w}_j = \mathbf{w}^T \mathbf{x}$

2. Update weights with the following definition: $\mathbf{w}_j := \mathbf{w}_j + \Delta \mathbf{w}_j$

Equals positive or negative value, so weights change

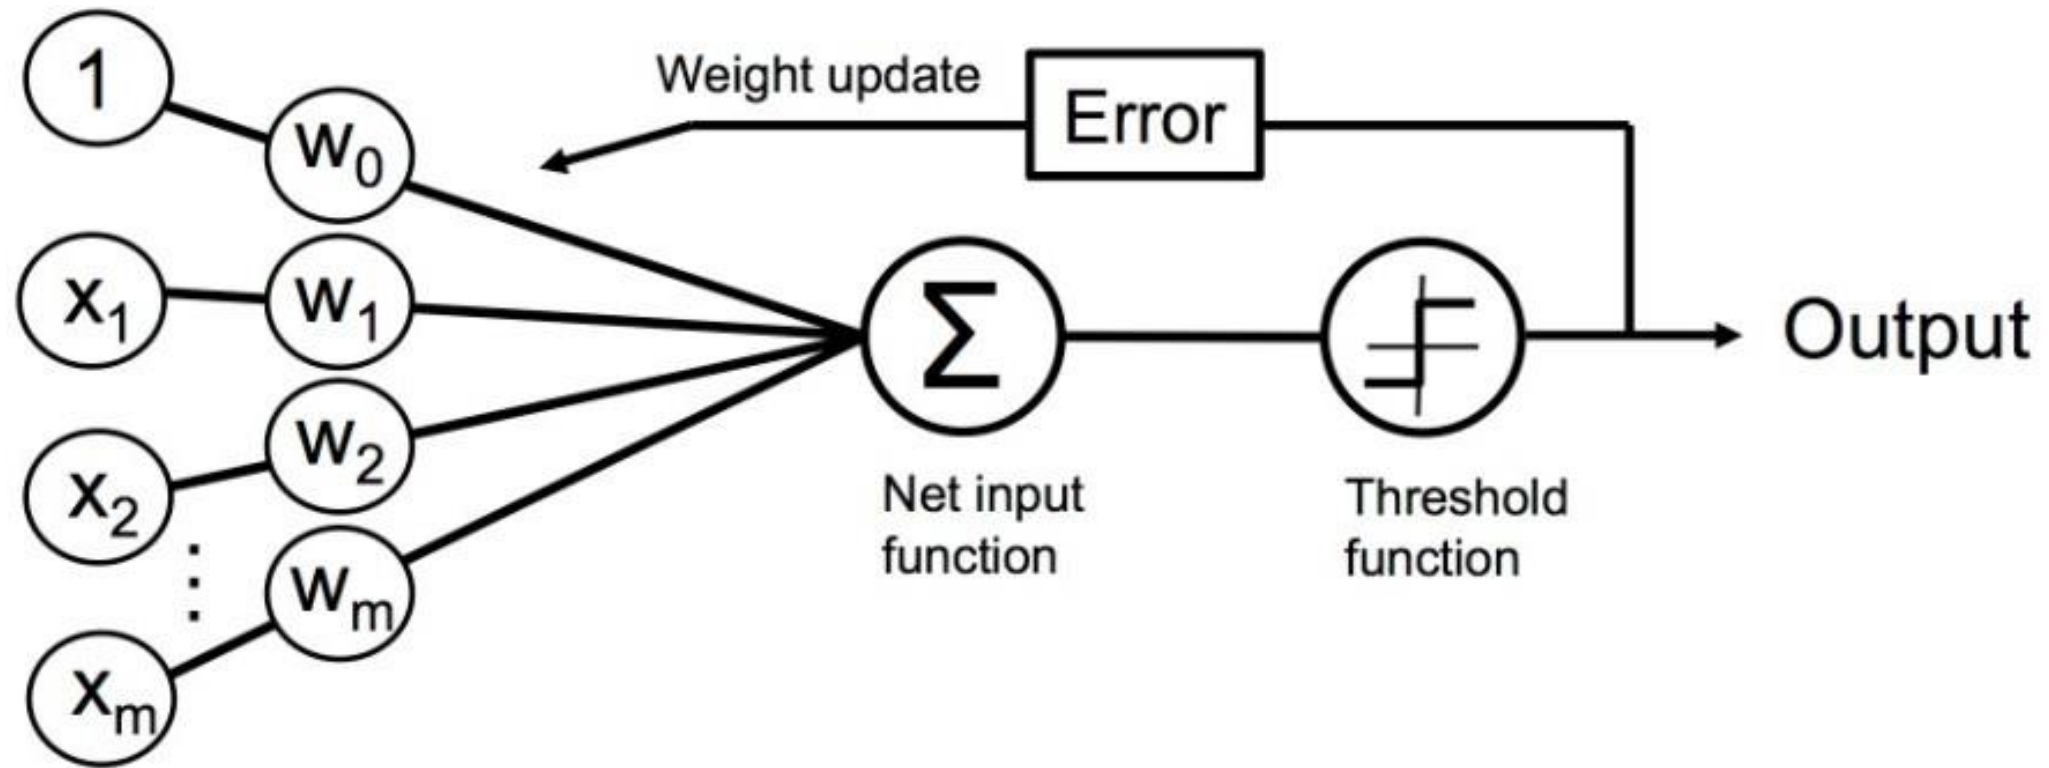
$$\Delta \mathbf{w}_j = \eta (\text{target}^{(i)} - \text{output}^{(i)}) \mathbf{x}_j^{(i)}$$

Learning Rate

True Class Label

Predicted Class Label

Perceptron: Learning Algorithm



Perceptron: Learning Algorithm Choices

- Learning rate
- Number of epochs (passes over the dataset)

Perceptron Limitation: XOR Problem


XOR = “Exclusive Or”

- Input: two binary values x_1 and x_2
- Output:
 - 1, when exactly one input equals 1
 - 0, otherwise

x_1	x_2	$x_1 \text{ XOR } x_2$
0	0	?
0	1	?
1	0	?
1	1	?

Perceptron Limitation: XOR Problem

Cannot solve XOR problem and so separate 1s from 0s with a perceptron (linear function):



x_1	x_2	$x_1 \text{ XOR } x_2$
0	0	0
0	1	1
1	0	1
1	1	0

Perceptron Limitation: XOR Problem




Frank Rosenblatt

How can a machine be “conscious”
when it can’t solve the XOR problem?


Neural Networks: Connected Neurons

Biological Neural Network:




<http://www.rzagabe.com/2014/11/03/an-introduction-to-artificial-neural-networks.html>

Artificial Neural Network:




Neural Network



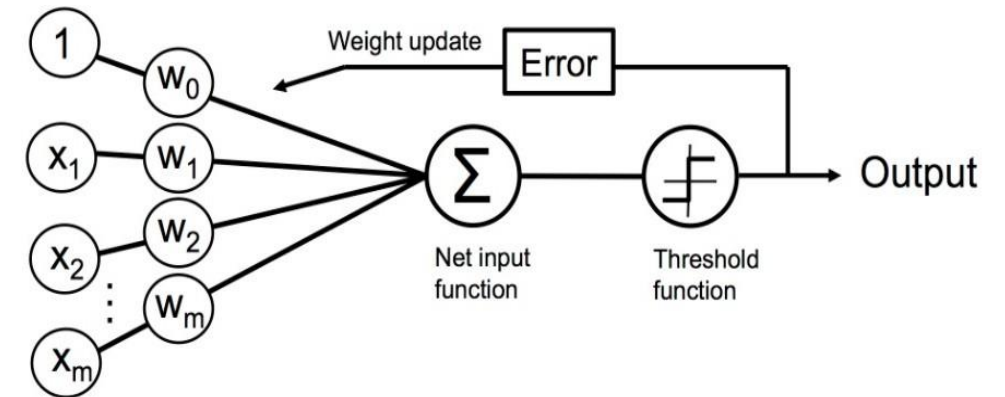
This is a 3-layer neural network (i.e., count number of hidden layers plus output layer)

each “hidden layer” uses outputs of units (i.e., neurons) and provides them as inputs to other units (i.e., neurons)

Neural Network



- How does this relate to a perceptron?




- Unit: takes as input a weighted sum and applies an activation function

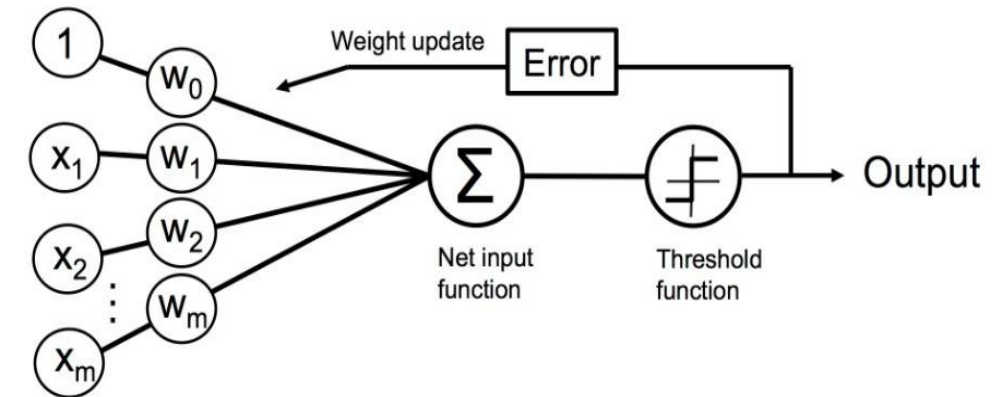
Python Machine Learning; Raschka & Mirjalili

[https://github.com/josephlee93/neural-networks-1/](https://github.com/josephlee93/neural-networks-1)

Neural Network



- How does this relate to a perceptron?




- Unit: takes as input a weighted sum and applies an activation function


Python Machine Learning; Raschka & Mirjalili

<http://cs231n.github.io/neural-networks-1/>

Neural Network



- How does this relate to a perceptron?




- Unit: takes as input a weighted sum and applies an activation function


Python Machine Learning; Raschka & Mirjalili

<http://cs231n.github.io/neural-networks-1/>

Neural Network



- How does this relate to a perceptron?




- Unit: takes as input a weighted sum and applies an activation function


Python Machine Learning; Raschka & Mirjalili

<http://cs231n.github.io/neural-networks-1/>

Neural Network



- How does this relate to a perceptron?




- Unit: takes as input a weighted sum and applies an activation function

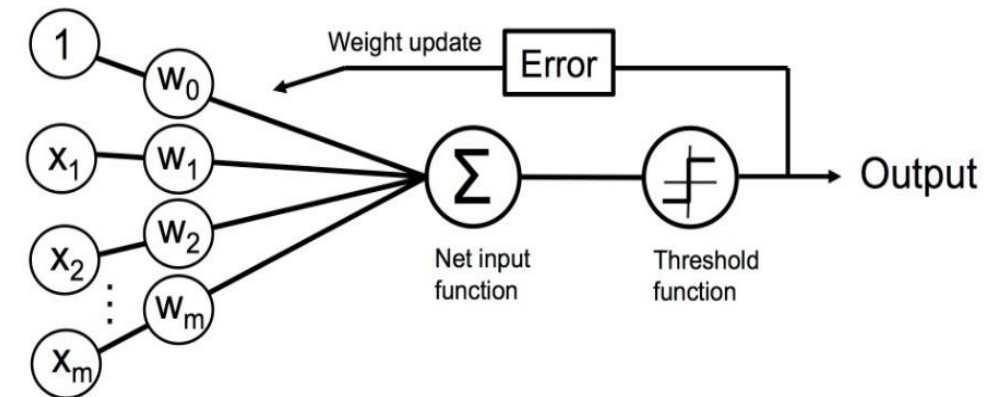
Python Machine Learning; Raschka & Mirjalili

<http://cs231n.github.io/neural-networks-1/>

Neural Network



- How does this relate to a perceptron?




- Unit: takes as input a weighted sum and applies an activation function


Python Machine Learning; Raschka & Mirjalili

<http://cs231n.github.io/neural-networks-1/>

Neural Network



- How does this relate to a perceptron?




- Unit: takes as input a weighted sum and applies an activation function


Python Machine Learning; Raschka & Mirjalili

<http://cs231n.github.io/neural-networks-1/>

Neural Network



- How does this relate to a perceptron?




- Unit: takes as input a weighted sum and applies an activation function

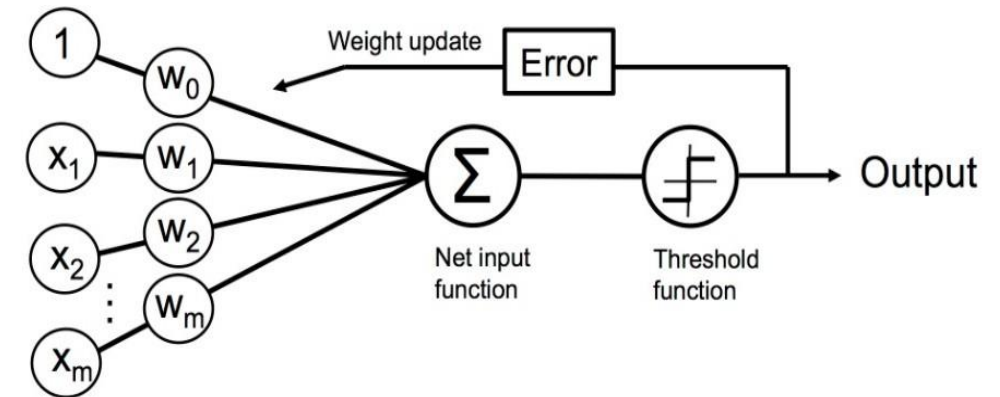
Python Machine Learning; Raschka & Mirjalili

<http://cs231n.github.io/neural-networks-1/>

Neural Network



- How does this relate to a perceptron?




- Unit: takes as input a weighted sum and applies an activation function

Python Machine Learning; Raschka & Mirjalili


<http://cs231n.github.io/neural-networks-1/>

Neural Network



- **Training goal: learn model parameters**
- Layers are called “hidden” because algorithm decides how to use each layer to produce its output


Neural Network



How many weights are in this model?

- Input to Hidden Layer 1:
 - $3 \times 4 = 12$
- Hidden Layer 1 to Hidden Layer 2:
 - $4 \times 4 = 16$
- Hidden Layer 2 to Output Layer
 - $4 \times 1 = 4$
- Total:
 - $12 + 16 + 4 = 32$


Neural Network



How many parameters are there to learn?

- Number of weights:
 - 32
- Number of biases:
 - $4 + 4 + 1 = 9$
- Total
 - 41

Fully Connected, Feedforward Neural Networks




- What does it mean for a model to be fully connected?
 - Each unit provides input to each unit in the next layer
- What does it mean for a model to be feedforward?
 - Each layer serves as input to the next layer with no loops

Hidden Layers Alone Are NOT Enough to Model Non-Linear Functions

Key Observation: feedforward networks are just functions chained together

e.g.,




- What is function for h_1 ?
 - $h_1 = w_1x_1 + w_3x_2 + b_1$
- What is function for h_2 ?
 - $h_2 = w_2x_1 + w_4x_2 + b_2$
- What is function for y ?
 - $y = h_1w_5 + h_2w_6 + b_3$
 - $y = (w_1x_1 + w_3x_2 + b_1)w_5 + (w_2x_1 + w_4x_2 + b_2)w_6 + b_3$
 - $y = w_1w_5x_1 + w_3w_5x_2 + w_5b_1 + w_2w_6x_1 + w_4w_6x_2 + w_6b_2 + b_3$

A chain of LINEAR functions at any depth is still a LINEAR function!

Hidden Layers Alone Are NOT Enough to Model Non-Linear Functions

Key Observation: feedforward networks are just functions chained together

e.g.,



- What is function for h_1 ?

- $h_1 = w_1x_1 + w_3x_2 + b_1$

- What is function for h_2 ?

- $h_2 = w_2x_1 + w_4x_2 + b_2$

- What is function for y ?

- $y = h_1w_5 + h_2w_6 + b_3$




Constant x linear function = linear function

A chain of LINEAR functions at any depth is still a LINEAR function!


Key Idea: Use Connected Neurons to Non-linearly Transform Input into Useful Features for Predictions

Biological Neural Network:




<http://www.rzagabe.com/2014/11/03/an-introduction-to-artificial-neural-networks.html>

Artificial Neural Network:




Key Idea: Use Connected Neurons to Non-linearly Transform Input into Useful Features for Predictions

Biological Neuron:




Artificial Neurons
(e.g., Perceptron):




Mimic a neuron firing,
by having each unit
apply a non-linear
“activation” function
to the weighted input

Non-Linear Activation Functions

- Each unit applies a non-linear “activation” function to the weighted input to mimic a neuron firing




$$\sigma(z) = \frac{1}{1 + \exp(-z)}$$



$$\tanh(z) = \frac{\exp(z) - \exp(-z)}{\exp(z) + \exp(-z)}$$


Computationally faster



$$\text{ReLU}(z) = \max(0, z)$$

Non-Linear Example: Revisiting XOR problem

- Non-linear function: separate 1s from 0s:




- Approach: ReLU activation function ($\text{ReLU}(z) = \max(0, z)$) with these parameters:


Neural networks can solve XOR problem...
and so model non-linear functions!

Activation Functions and Model Parameters (e.g., Sigmoid)

Weights determine curvature:




Biases determine shifted position:




Desired Output Driven by Task


Regression
(predict **continuous** value)



Classification
(predict **discrete** value)



Sigmoid (for Binary Classification)




$$\sigma(z) = \frac{1}{1 + \exp(-z)}$$

If ≥ 0.5 , output 1;

Else, outputs 0

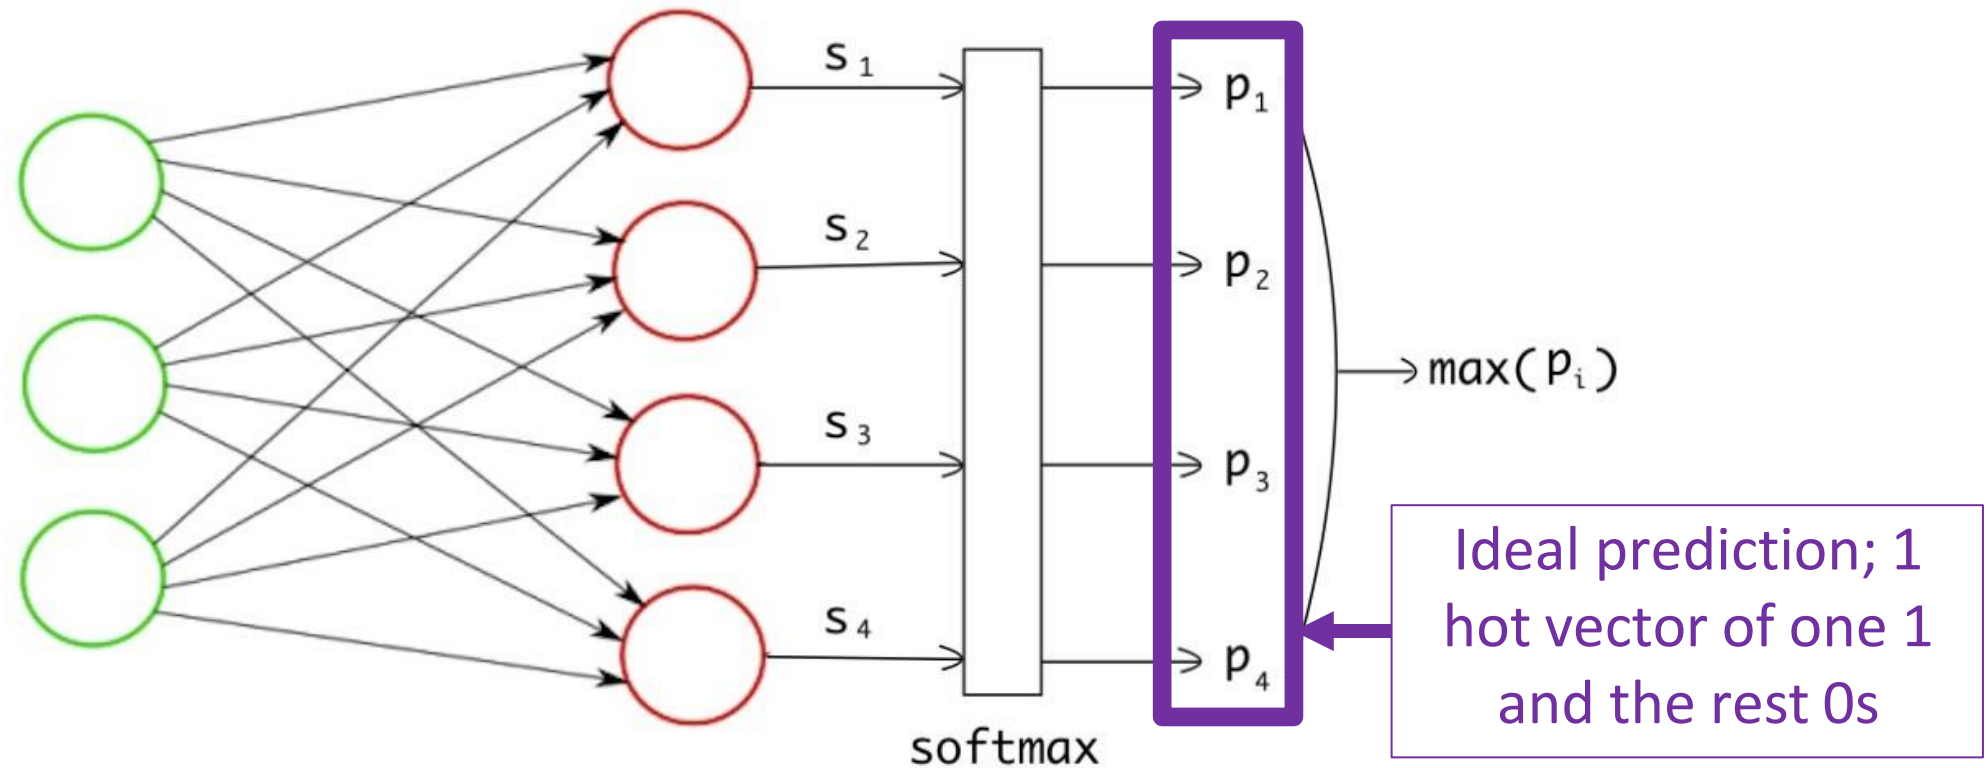
Softmax (for Multiclass Classification)

Converts vector of **scores** into a **probability distribution** that sums to 1; e.g.,



Softmax (for Multiclass Classification)

Converts vector of **scores** into a **probability distribution** that sums to 1; e.g.,



Softmax (for Multiclass Classification)

Converts vector of **scores** into a probability distribution that sums to 1

Get rid of negative values while preserving original order of scores; e causes negative scores to become slightly larger than 0 while positive values grow exponentially (choosing e rather than another exponent base simplifies math during training)

$$\sigma(\mathbf{z})_i = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$


$i = 1, \dots, K$

Number of classes

Want to divide each node's score by sum of all entries to make them sum to 1 (normalization)

Softmax (for Multiclass Classification)

Converts vector of **scores** into a probability distribution that sums to 1; e.g.,




Softmax (for Multiclass Classification)

Converts vector of **scores** into a probability distribution that sums to 1; e.g.,

	Scoring Function
Dog	-3.44
Cat	1.16
Boat	-0.81
Airplane	3.91

Softmax (for Multiclass Classification)

Converts vector of scores into a **probability distribution** that sums to 1; e.g.,



Softmax (for Multiclass Classification)


Converts vector of scores into a **probability distribution** that sums to 1; e.g.,

$$e^{z_i} \quad \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$

	Scoring Function	Unnormalized Probabilities	Normalized Probabilities
Dog	-3.44	0.0321	0.0006
Cat	1.16	3.1899	0.0596
Boat	-0.81	0.4449	0.0083
Airplane	3.91	49.8990	0.9315

Softmax (for Multiclass Classification)

Converts vector of **scores** into a **probability distribution** that sums to 1; e.g.,




Softmax (for Multiclass Classification)

Converts vector of scores into a probability distribution that sums to 1; e.g.,


	Scoring Function	Unnormalized Probabilities	Normalized Probabilities
Dog	-3.44	0.0321	0.0006
Cat	1.16	3.1899	0.0596
Boat	-0.81	0.4449	0.0083
Airplane	3.91	49.8990	0.9315

Desired Output Driven by Task

Regression
(predict **continuous** value)



Classification
(predict **discrete** value)



Desired Output Driven by Task

