

# SEQUENCE MODELS

Dr Usman Zia



## Motivation: Need for Sequential Modeling

 Why do we need *Sequential Modeling*?

# Motivation: Need for Sequential Modeling

## Examples of Sequence data

- **Speech Recognition**



*This is RNN*

- **Machine Translation**

Hello, I am usman.

Hallo, ich bin Usman

- **Language Modeling**

Recurrent neural ? based ? model

network

language

- **Named Entity Recognition**

David lives in Munich

David lives in Munich  
*person location*

- **Sentiment Classification**

There is nothing to like in this movie.



- **Video Activity Analysis**



Punching

# Motivation: Need for Sequential Modeling

Inputs, Outputs can be different lengths in different examples

*Example:*

Sentence1: David lives in Munich

Sentence2: David Clark lives in Munich DE

# Motivation: Need for Sequential Modeling

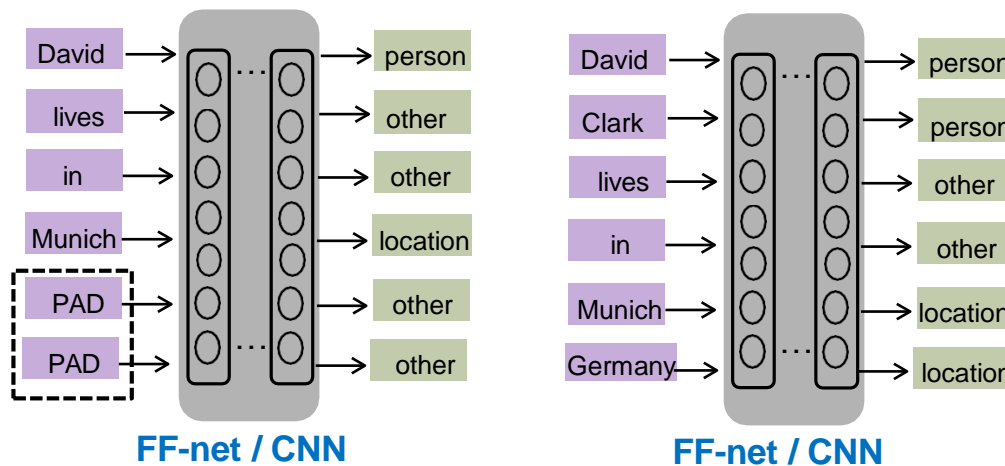
Inputs, Outputs can be different lengths in different examples

Example:

Sentence1: David lives in Munich

Sentence2: David Clark lives in Munich DE

Additional word  
'PAD' i.e., padding



\*FF-net: Feed-forward network


# Motivation: Need for Sequential Modeling

Inputs, Outputs can be different lengths in different examples


Example:

Sentence1: David lives in Munich

Sentence2: David Clark lives in Munich DE




FF-net / CNN




FF-net / CNN

\*FF-net: Feed-forward network



Models variable length sequences



Sequential model: RNN

# Motivation: Need for Sequential Modeling

Share Features learned across different positions or time steps

*Example:*

Sentence1: *Market falls into bear territory* → *Trading/Marketing*

Sentence2: *Bear falls into market territory* → *UNK*

**Same uni-gram  
statistics**



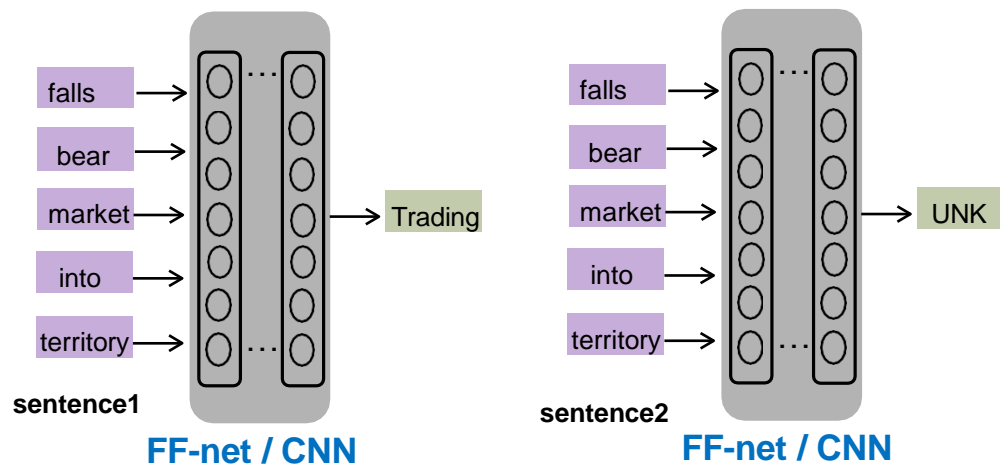
# Motivation: Need for Sequential Modeling

Share Features learned across different positions or time steps

Example:

Sentence1: *Market falls into bear territory* → *Trading/Marketing*

Sentence2: *Bear falls into market territory* → *UNK*



No sequential  
or temporal  
modeling, i.e.,  
order-less

Treats the two  
sentences the  
same




# Motivation: Need for Sequential Modeling

Share Features learned across different positions or time steps

Example:

Sentence1: *Market falls into bear territory* → *Trading/Marketing*

Sentence2: *Bear falls into market territory* → *UNK*



Language concepts,  
Word ordering,  
Syntactic & semantic information

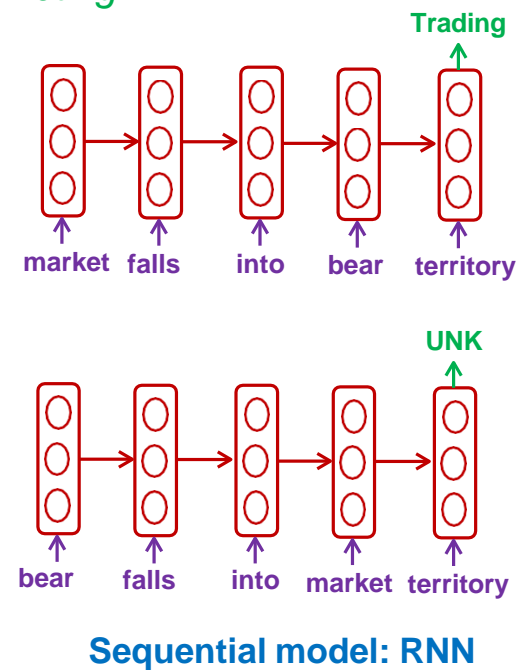

# Motivation: Need for Sequential Modeling

Share Features learned across different positions or time steps

Example:

Sentence1: *Market falls into bear territory* → *Trading/Marketing*


Sentence2: *Bear falls into market territory* → *UNK*



Language concepts,  
Word ordering,  
Syntactic & semantic information

# Motivation: Need for Sequential Modeling

**Machine Translation:** Different Input and Output sizes, incurring sequential patterns



# Long Term and Short Dependencies

## Short Term Dependencies

→ need recent information to perform the present task.

For example in a language model, predict the next word based on the previous ones.

*“the clouds are in the ?”* → ‘sky’ *“the clouds are in the **sky**”*

→ Easier to predict ‘sky’ given the context, i.e., *short term dependency*

## Long Term Dependencies

→ Consider longer word sequence “I grew up in France..... I speak fluent **French**.”

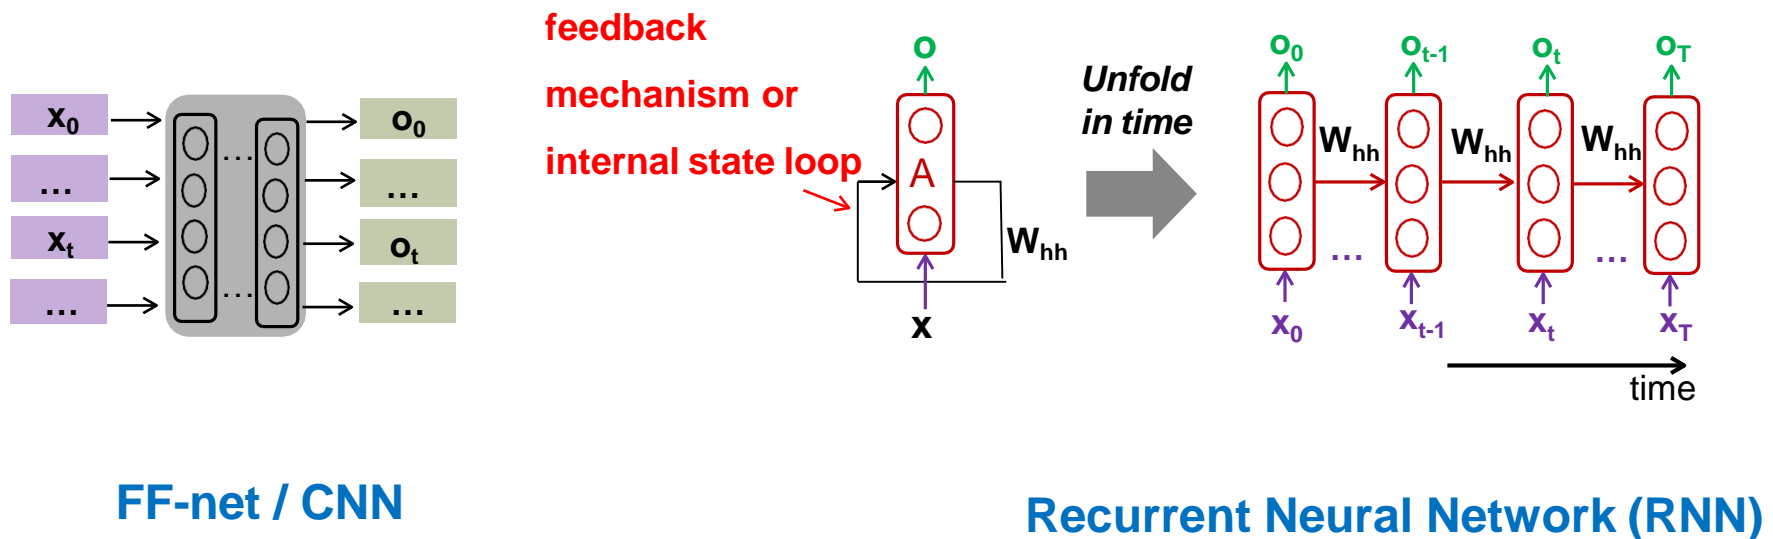
→ Recent information suggests that the next word is probably the name of a language, but if we want to narrow down which language, we need the context of France, from further back.

# Foundation of Recurrent Neural Networks

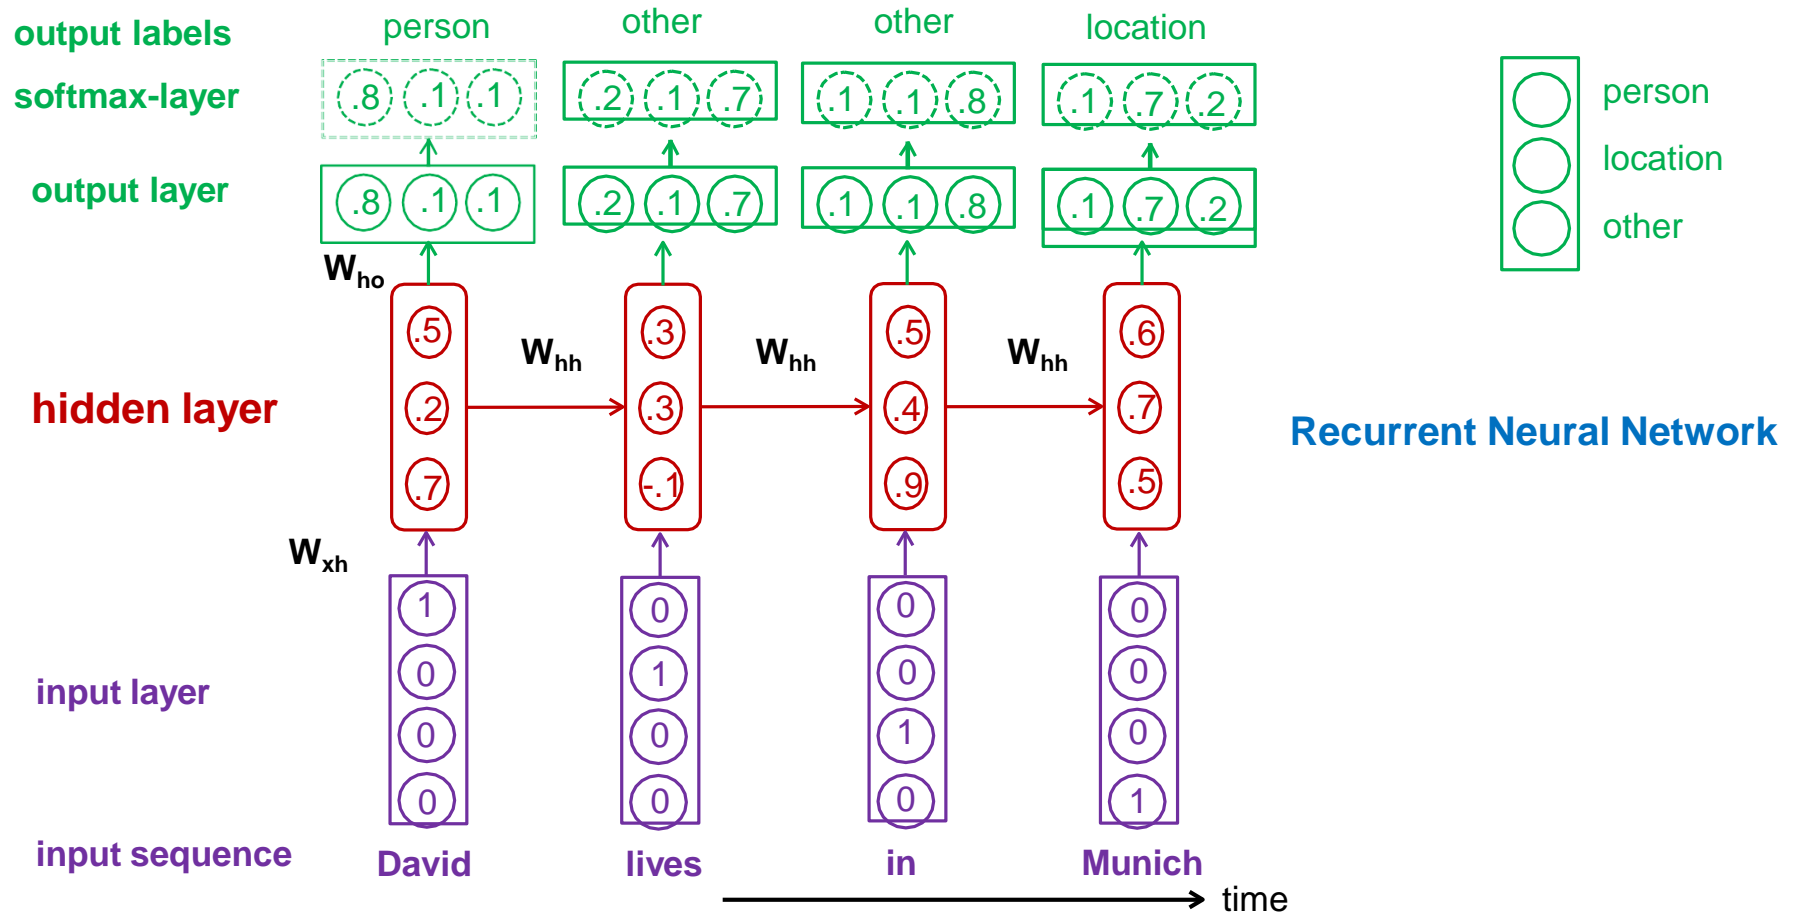
## Goal

- model long term dependencies
- connect previous information to the present task
- model sequence of events with loops, allowing information to persist


**Feed Forward NNets can not take time dependencies into account. Sequential data needs a Feedback Mechanism.**



# Foundation of Recurrent Neural Networks



# Recurrent Neural Networks (RNNs)





# Recurrent Neural Networks (RNNs)

We can process a sequence of vectors  $\mathbf{x}$  by applying a **recurrence formula** at every time step:


$$\boxed{h_t} = \boxed{f_W}(\boxed{h_{t-1}}, \boxed{x_t})$$

new state

some function with parameters  $W$

old state


input vector at some time step



Notice: the same function and the same set of parameters are used at every time step.

# Recurrent Neural Networks (RNNs)

The state consists of a single “*hidden*” vector  $\mathbf{h}$ :



$$h_t = f_W(h_{t-1}, x_t)$$





$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$


$$y_t = W_{hy}h_t$$

# RNN Computational Graph


Initial hidden state  
Either set to all 0,  
Or learn it



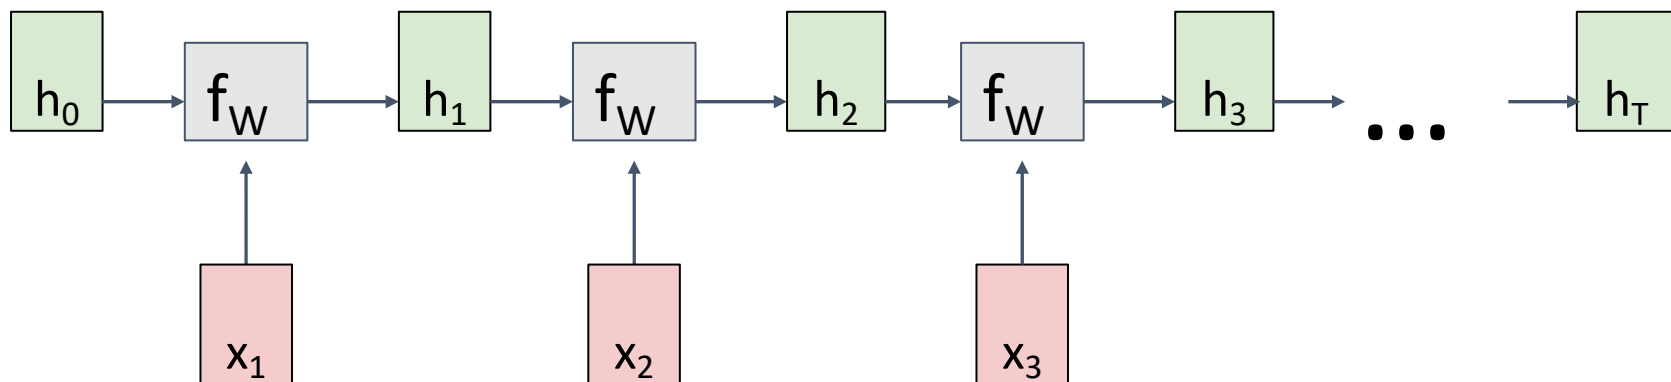
# RNN Computational Graph



# RNN Computational Graph

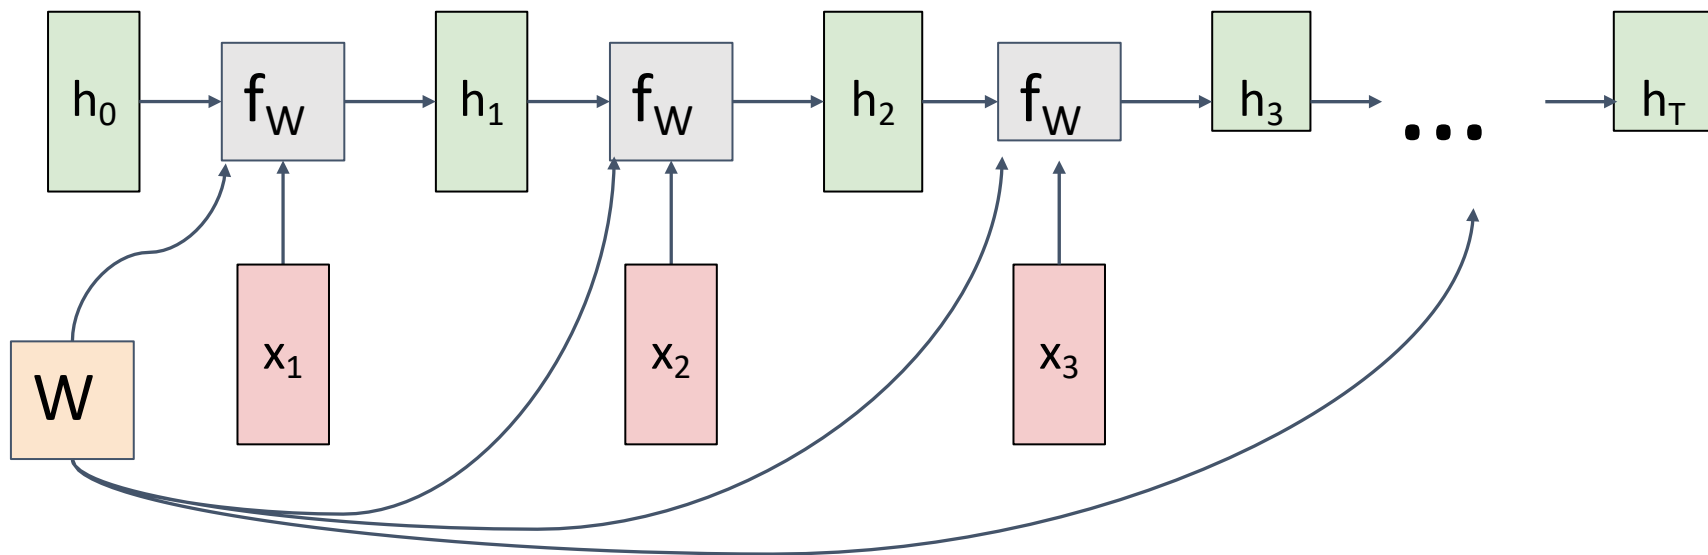


# RNN Computational Graph

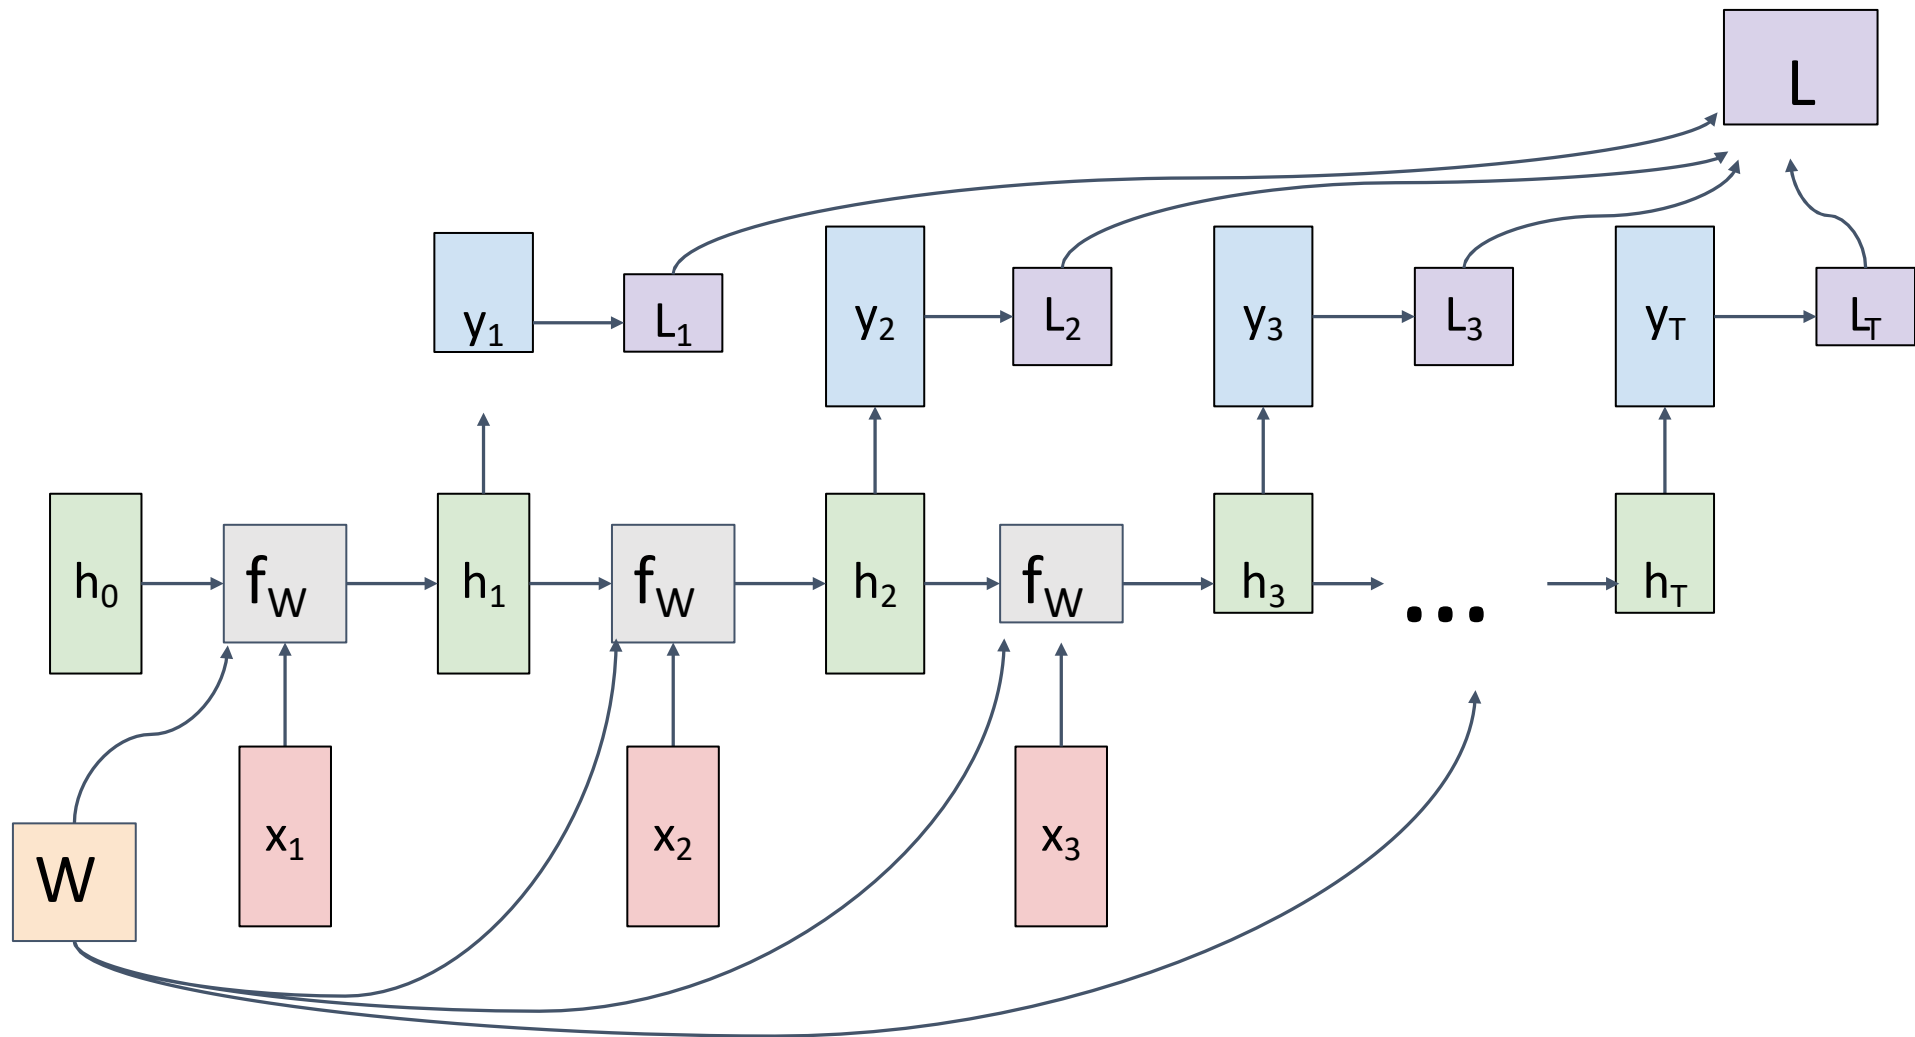


# RNN Computational Graph

Re-use the same weight matrix at every time-step

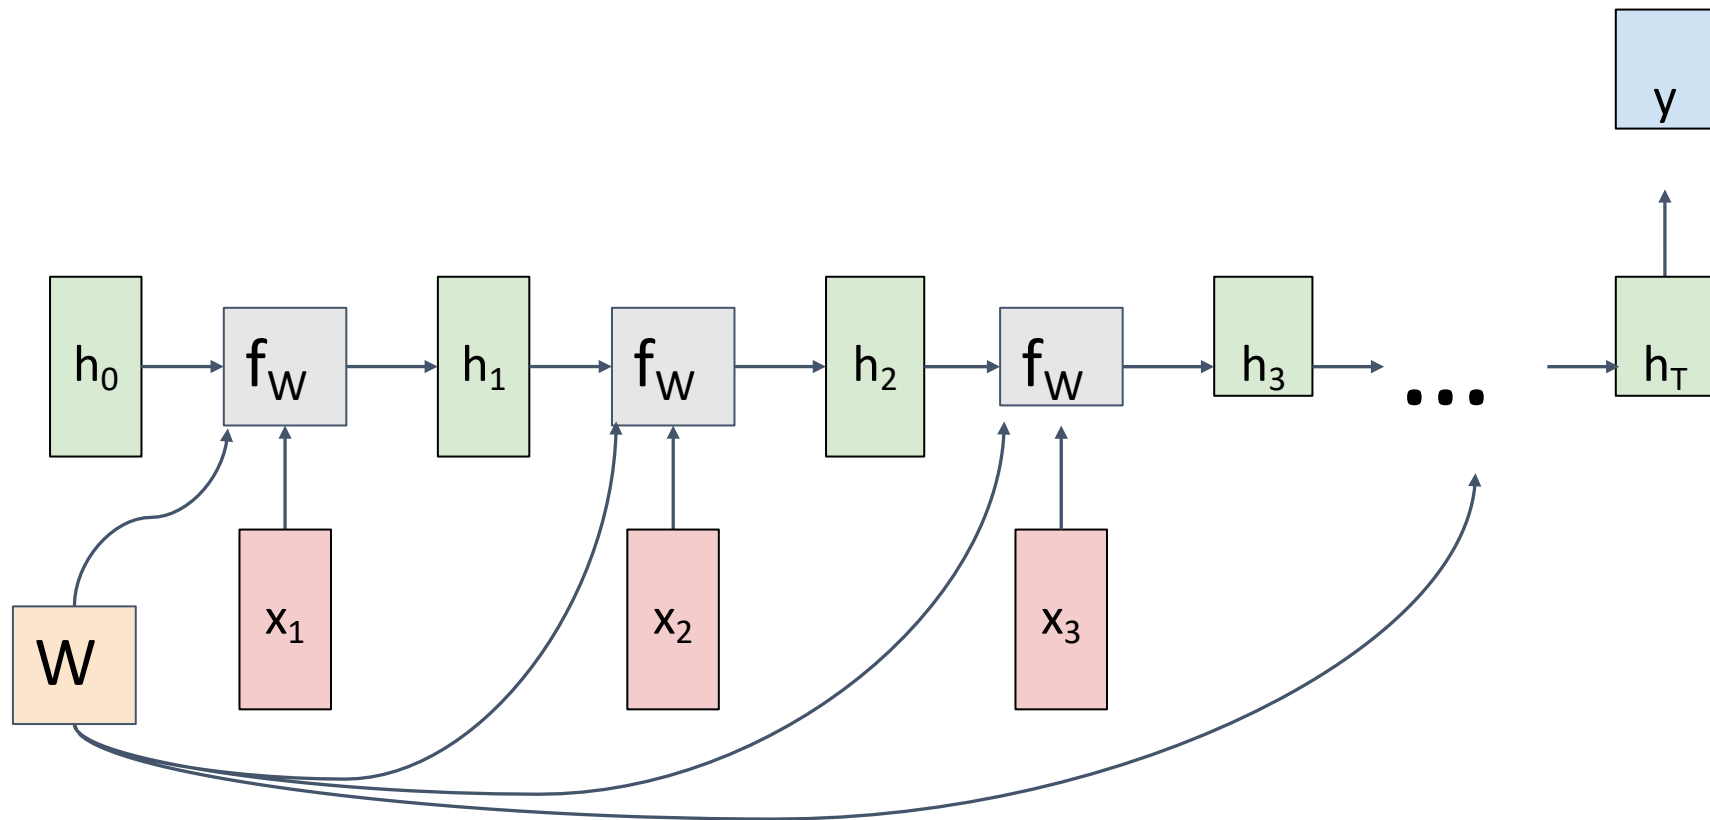


# RNN Computational Graph (Many to Many)

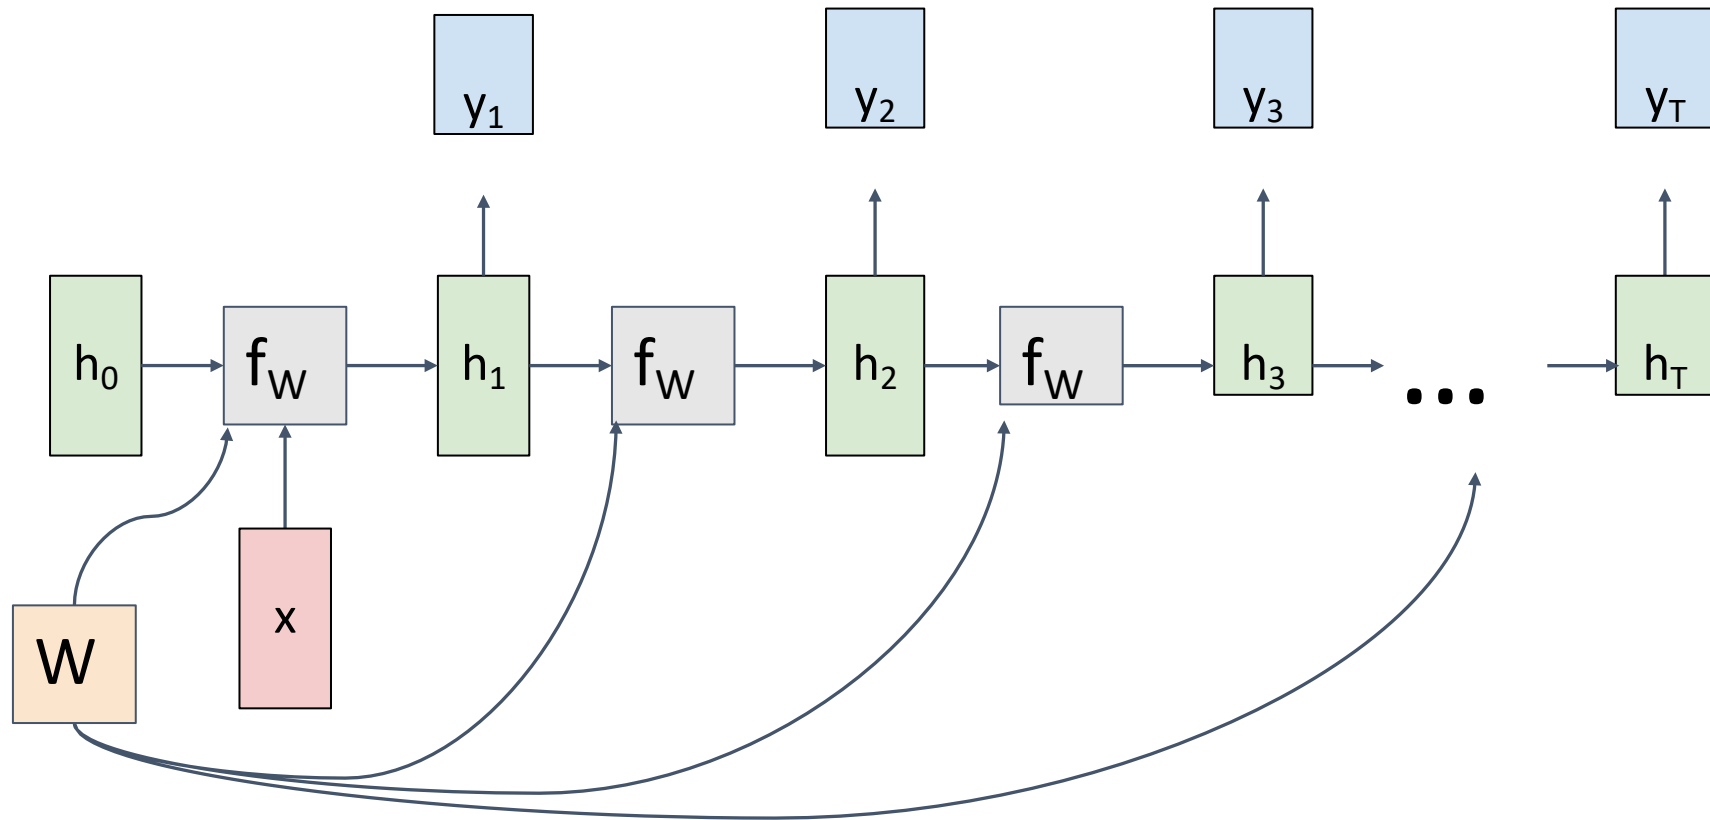




# RNN Computational Graph (Many to One)

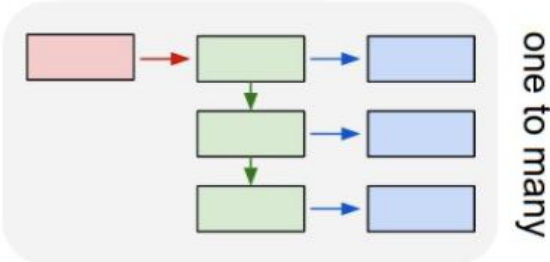

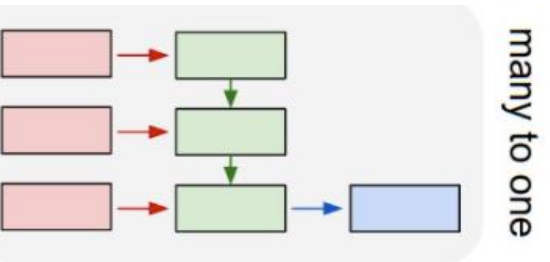
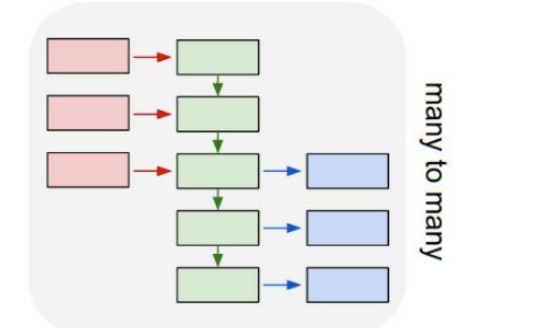


# RNN Computational Graph (One to Many)



# Recurrent Neural Networks (RNNs)

- Use cases for RNNs


RNN	Application	Input	Output
	Image Captioning		A person riding a motorbike on dirt road
	Sentiment Analysis	Awesome movie. Highly recommended.	Positive
	Machine Translation Video Classification	Happy Eid	عيد مبارك

# Example: Language Modeling

Given characters 1, 2, ..., t,  
model predicts character t

Training sequence: "hello"

Vocabulary: [h, e, l, o]



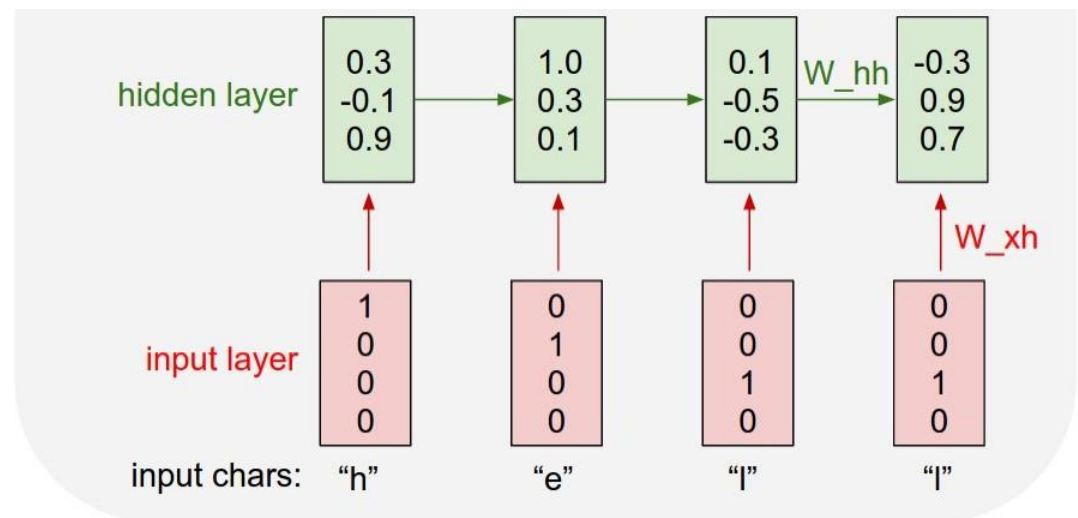
# Example: Language Modeling

Given characters 1, 2, ..., t,  
model predicts character t

$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

Training sequence: "hello"

Vocabulary: [h, e, l, o]



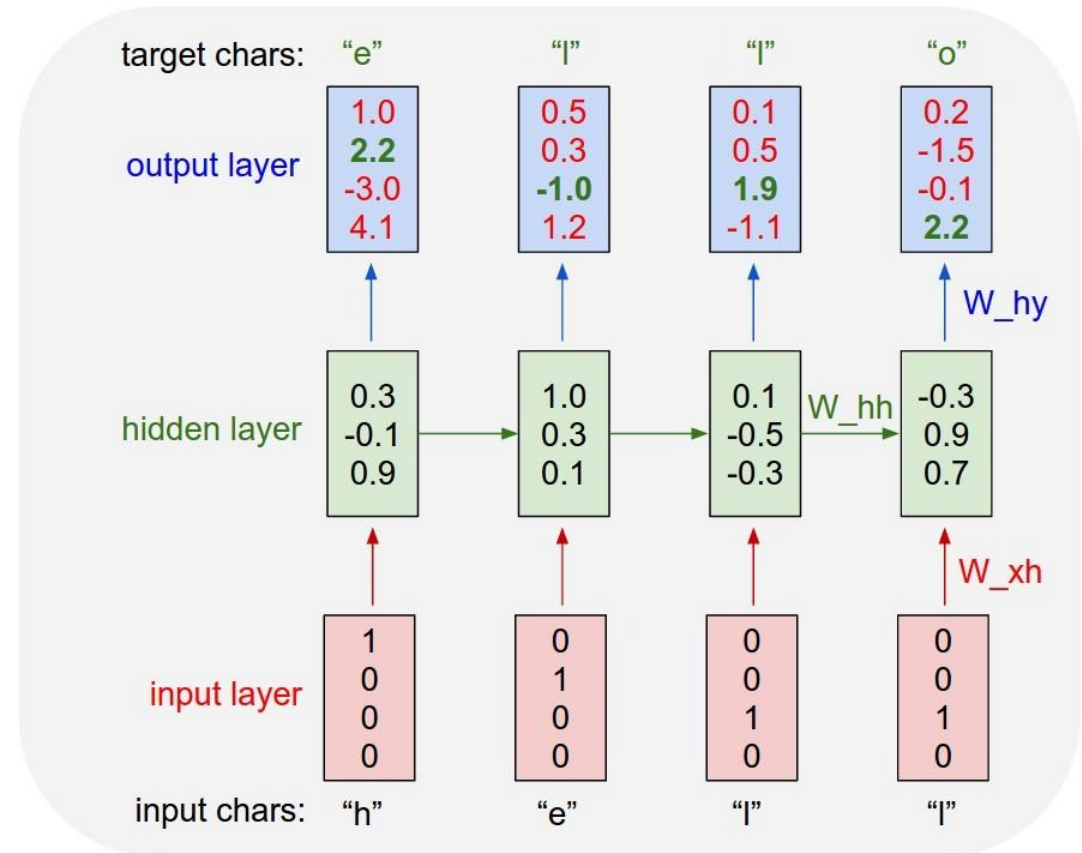
# Example: Language Modeling

Given characters 1, 2, ..., t,  
model predicts character t

$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

Training sequence: "hello"

Vocabulary: [h, e, l, o]



# Example: Language Modeling

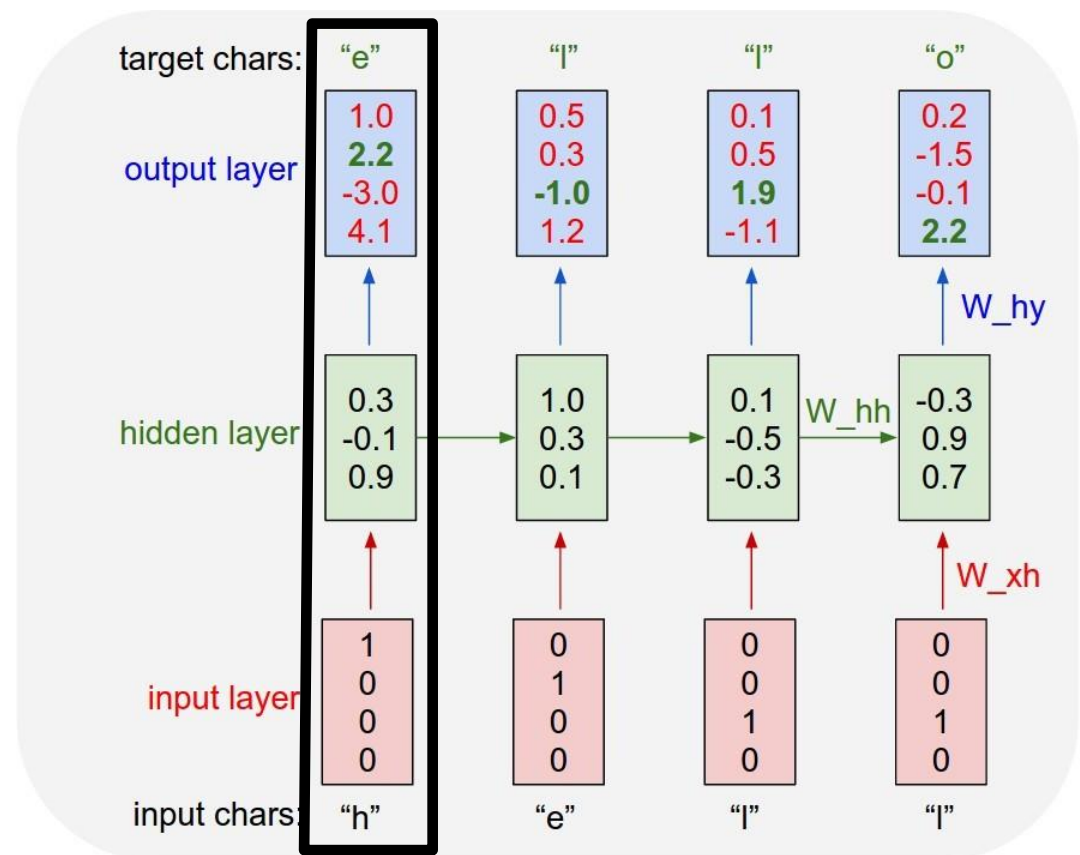
Given characters 1, 2, ..., t,  
model predicts character t

$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

Training sequence: "hello"

Vocabulary: [h, e, l, o]

Given "h", predict "e"



# Example: Language Modeling

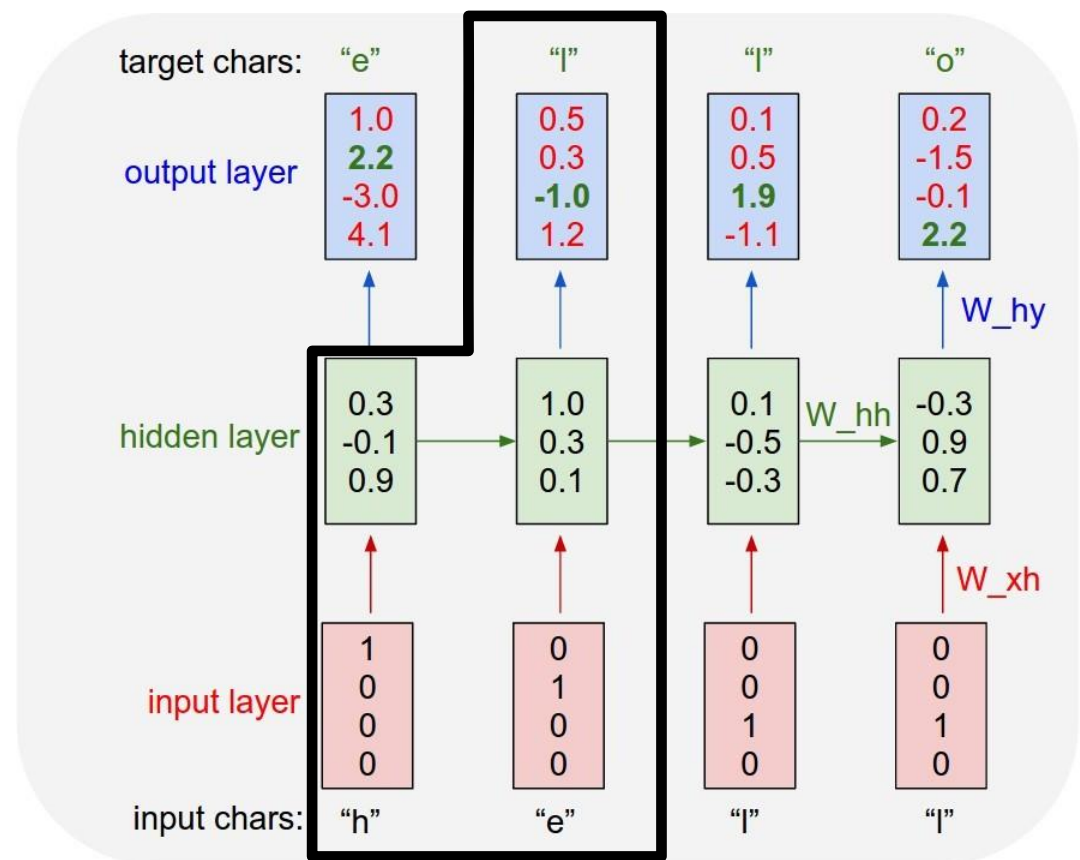
Given characters 1, 2, ..., t,  
model predicts character t

$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

Training sequence: "hello"

Vocabulary: [h, e, l, o]

Given "he", predict "l"





# Example: Language Modeling

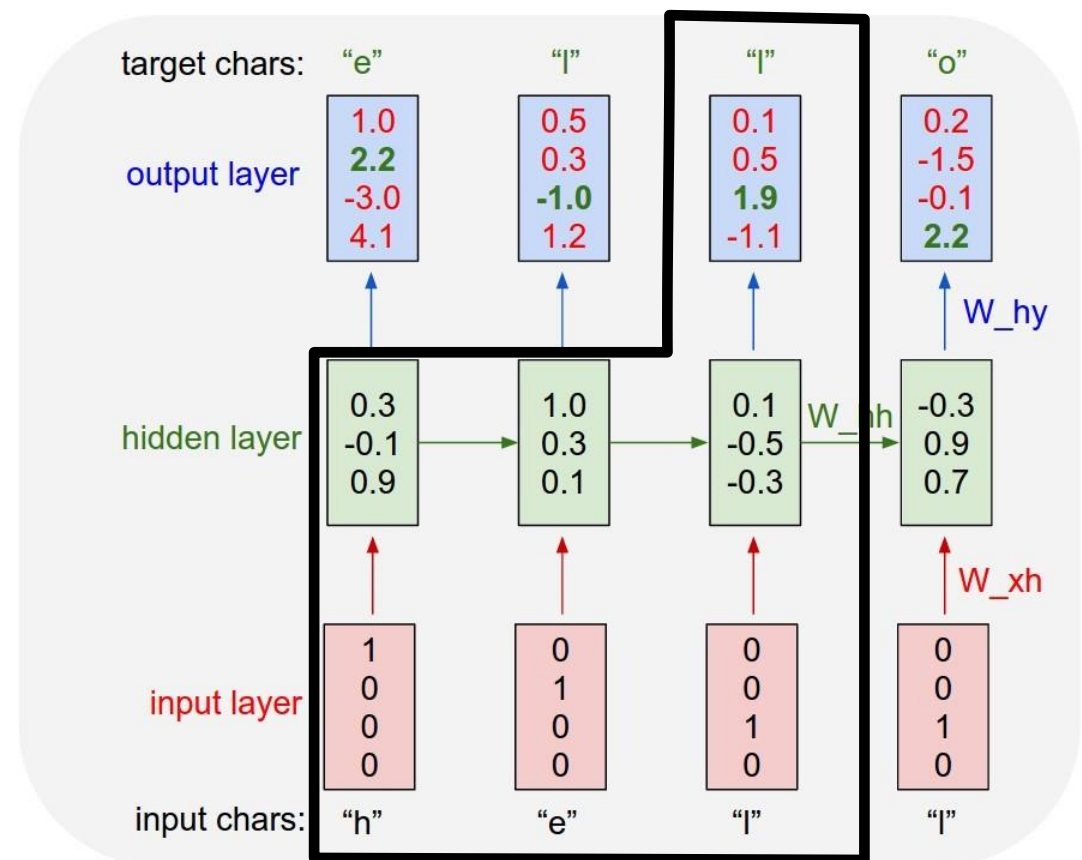
Given characters 1, 2, ..., t,  
model predicts character t

$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

Training sequence: "hello"

Vocabulary: [h, e, l, o]

Given "hel", predict "l"



# Example: Language Modeling

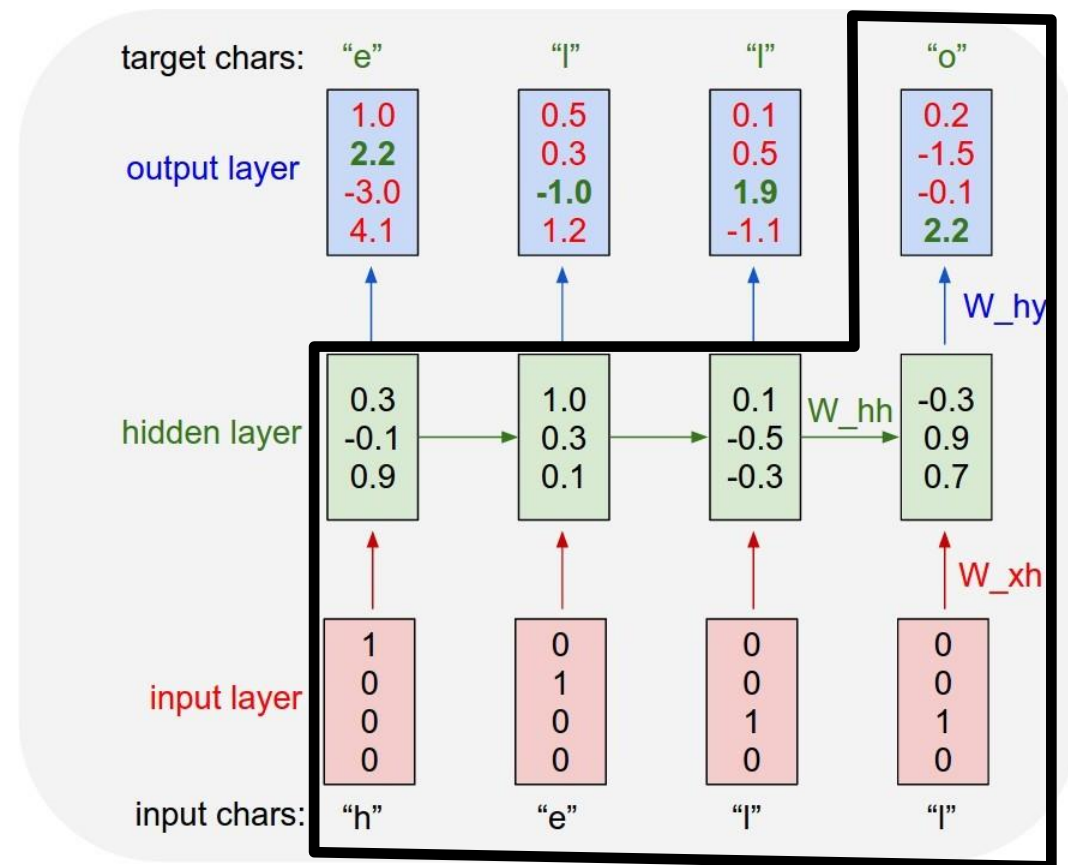
Given characters 1, 2, ..., t,  
model predicts character t

$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

Training sequence: "hello"

Vocabulary: [h, e, l, o]

Given "hell", predict "o"




# Example: Language Modeling

At test-time ... ???

Training sequence: "hello"

Vocabulary: [h, e, l, o]

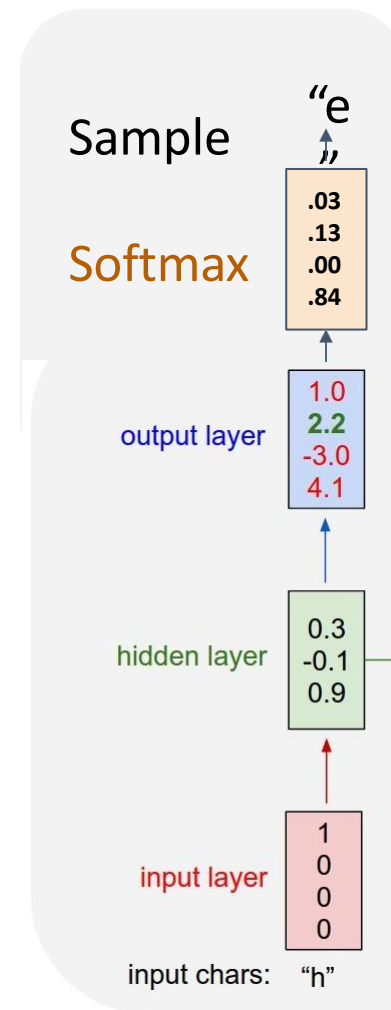


# Example: Language Modeling

At test-time, **generate** new text: sample characters one at a time, feed back to model

Training sequence: "hello"

Vocabulary: [h, e, l, o]




# Example: Language Modeling

At test-time, **generate** new text: sample characters one at a time, feed back to model

Training sequence: "hello"

Vocabulary: [h, e, l, o]

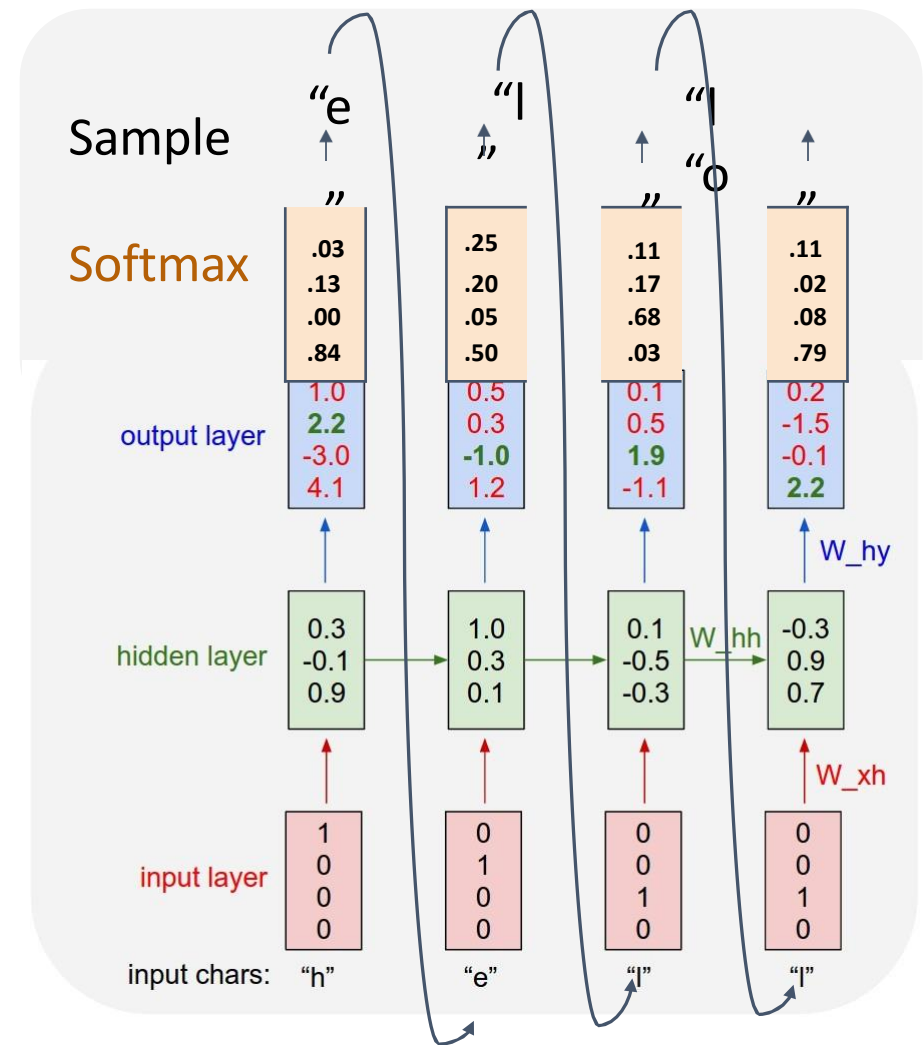


# Example: Language Modeling

At test-time, **generate** new text: sample characters one at a time, feed back to model

Training sequence: "hello"

Vocabulary: [h, e, l, o]

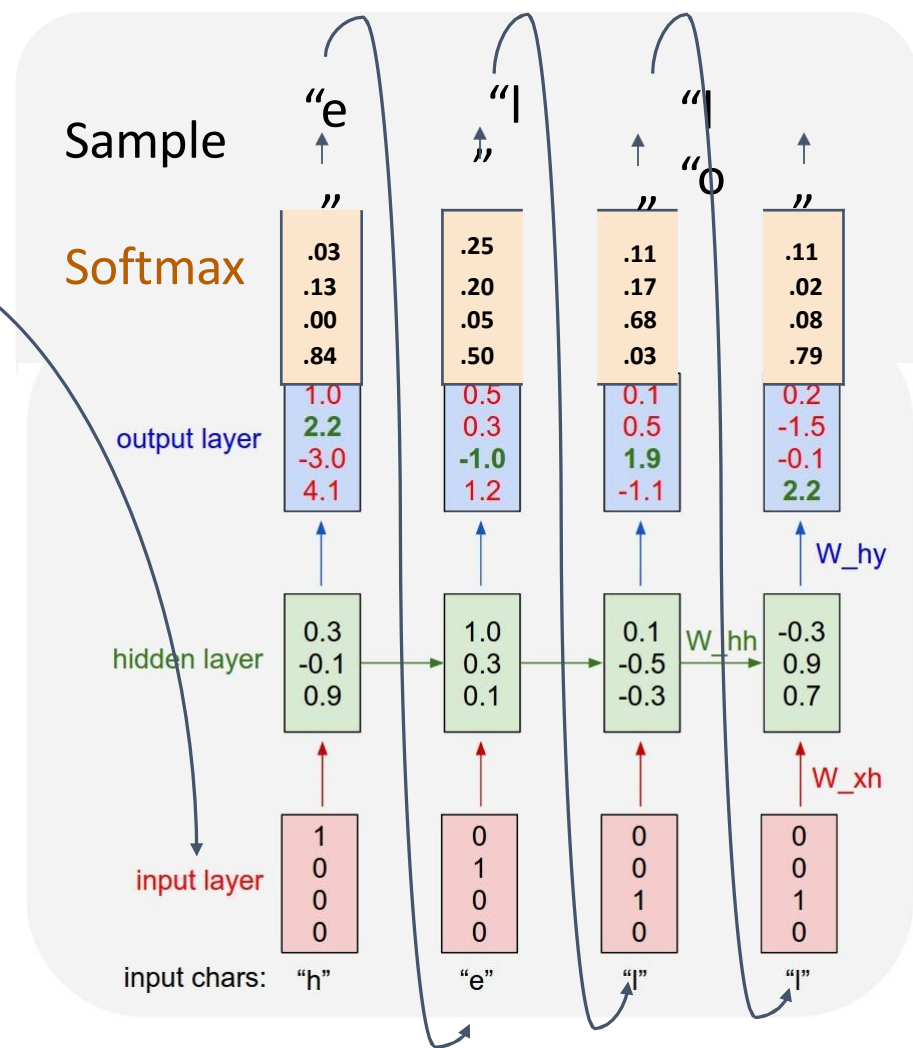


# Example: Language Modeling

So far: encode inputs as **one-hot-vector**

$$\begin{bmatrix} w_{11} & w_{12} & w_{13} & w_{14} \\ w_{21} & w_{22} & w_{23} & w_{24} \\ w_{31} & w_{32} & w_{33} & w_{34} \end{bmatrix} \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \end{bmatrix} = \begin{bmatrix} w_{11} \\ w_{21} \\ w_{31} \end{bmatrix}$$

Matrix multiply with a one-hot vector just extracts a column from the weight matrix. Often extract this into a separate **embedding** layer

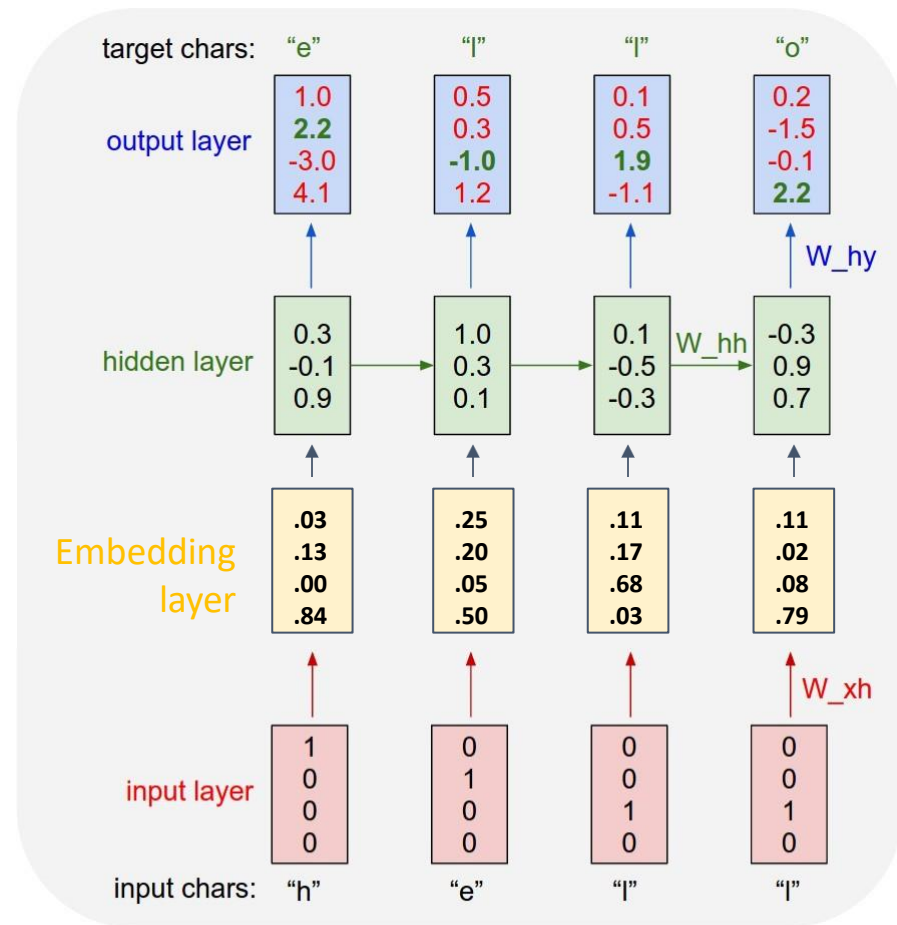


# Example: Language Modeling

So far: encode inputs as **one-hot-vector**



$$\begin{bmatrix} w_{11} & w_{12} & w_{13} & w_{14} \\ w_{21} & w_{22} & w_{23} & w_{24} \\ w_{31} & w_{32} & w_{33} & w_{34} \end{bmatrix} \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \end{bmatrix} = \begin{bmatrix} w_{11} \\ w_{21} \\ w_{31} \end{bmatrix}$$

Matrix multiply with a one-hot vector just extracts a column from the weight matrix. Often extract this into a separate **embedding** layer






# Example: Image Captioning



**Recurrent  
Neural  
Network**

**Convolutional Neural Network**

# Example: Image Captioning

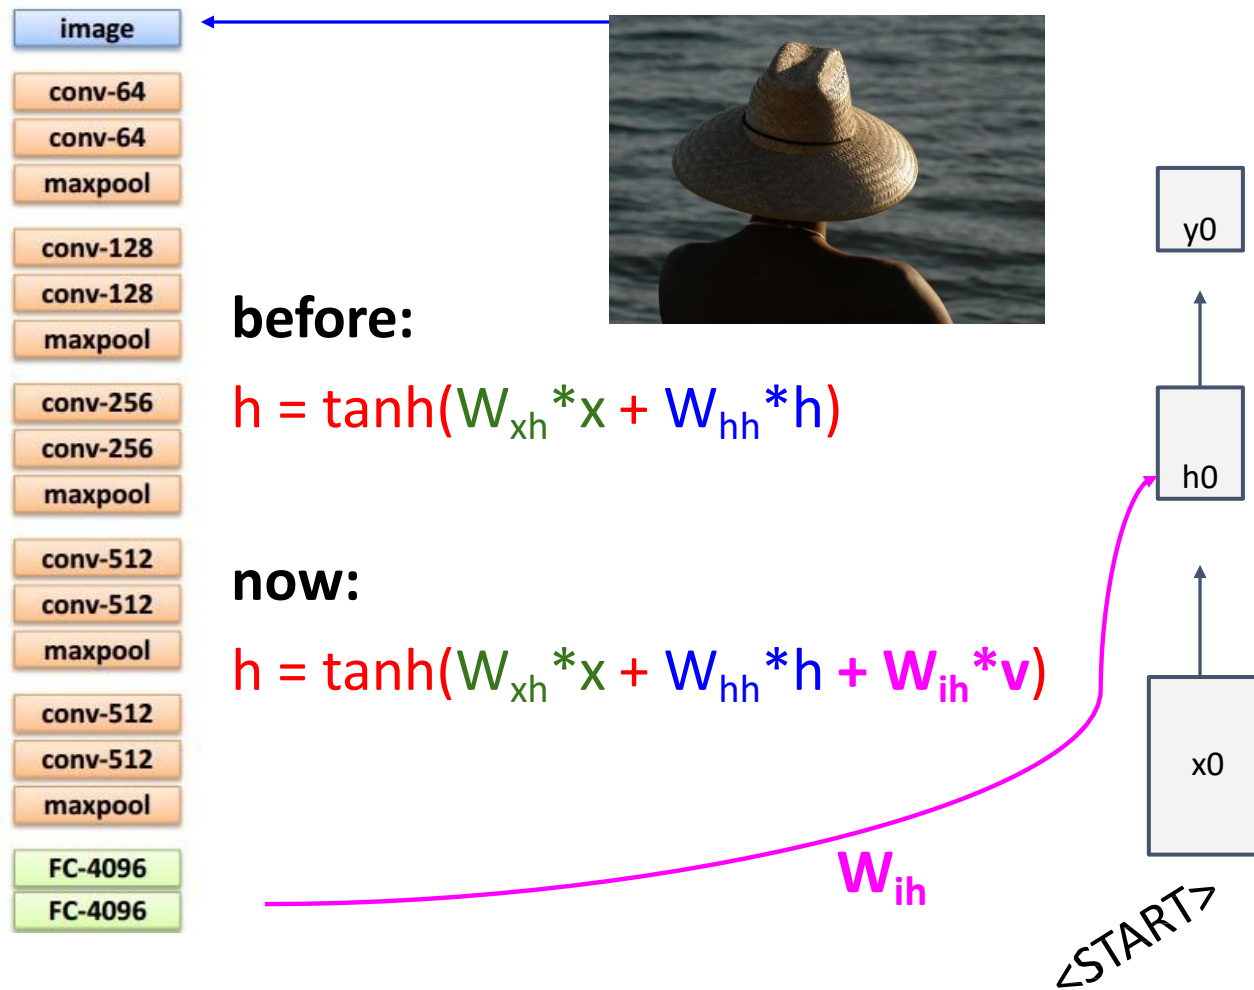


**Transfer learning:** Take CNN trained on ImageNet, chop off last layer


# Example: Image Captioning



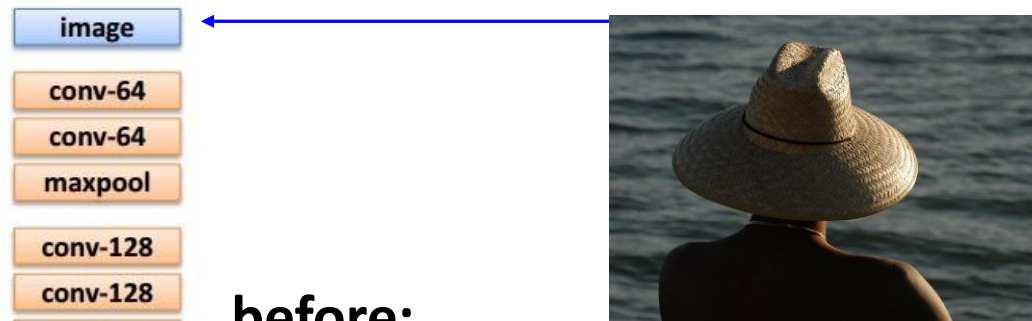
# Example: Image Captioning



# Example: Image Captioning



# Example: Image Captioning



**before:**

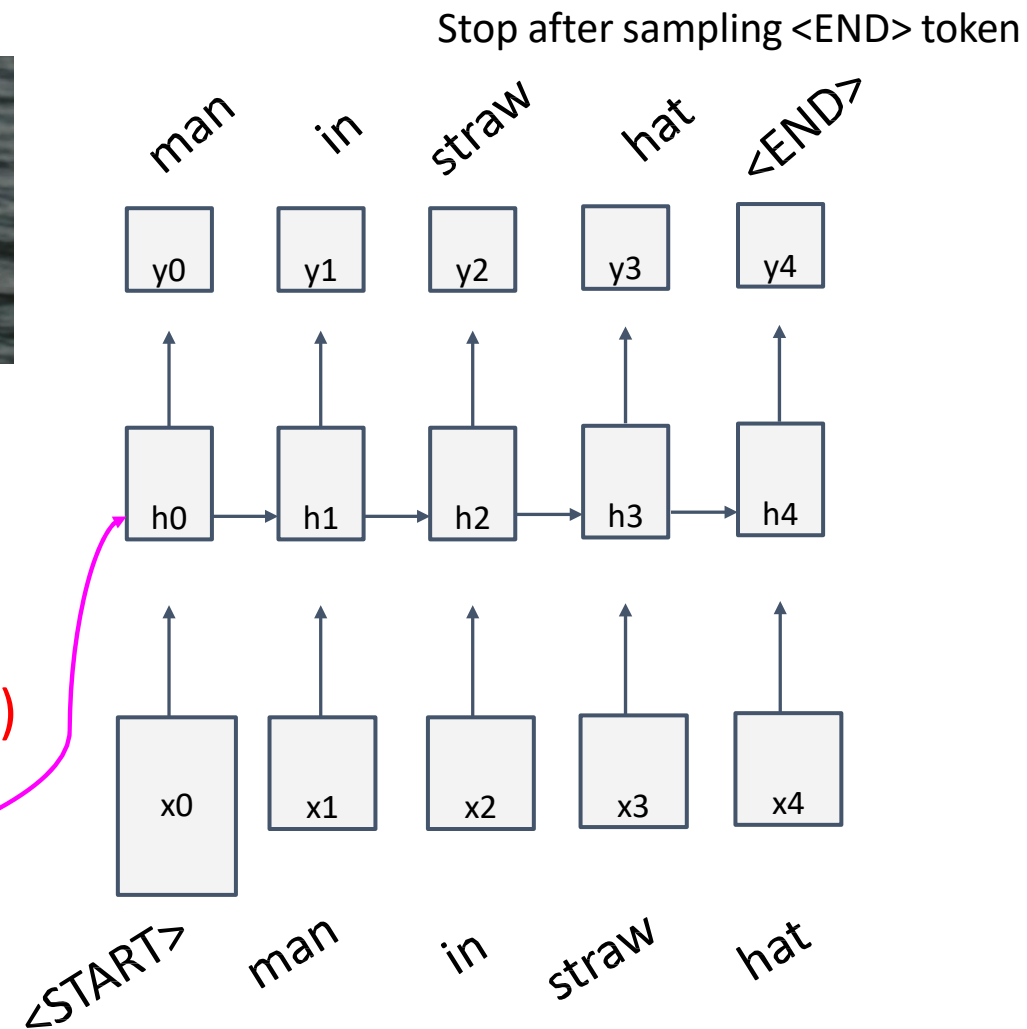
$$h = \tanh(W_{xh} * x + W_{hh} * h)$$

**now:**

$$h = \tanh(W_{xh} * x + W_{hh} * h + W_{ih} * v)$$



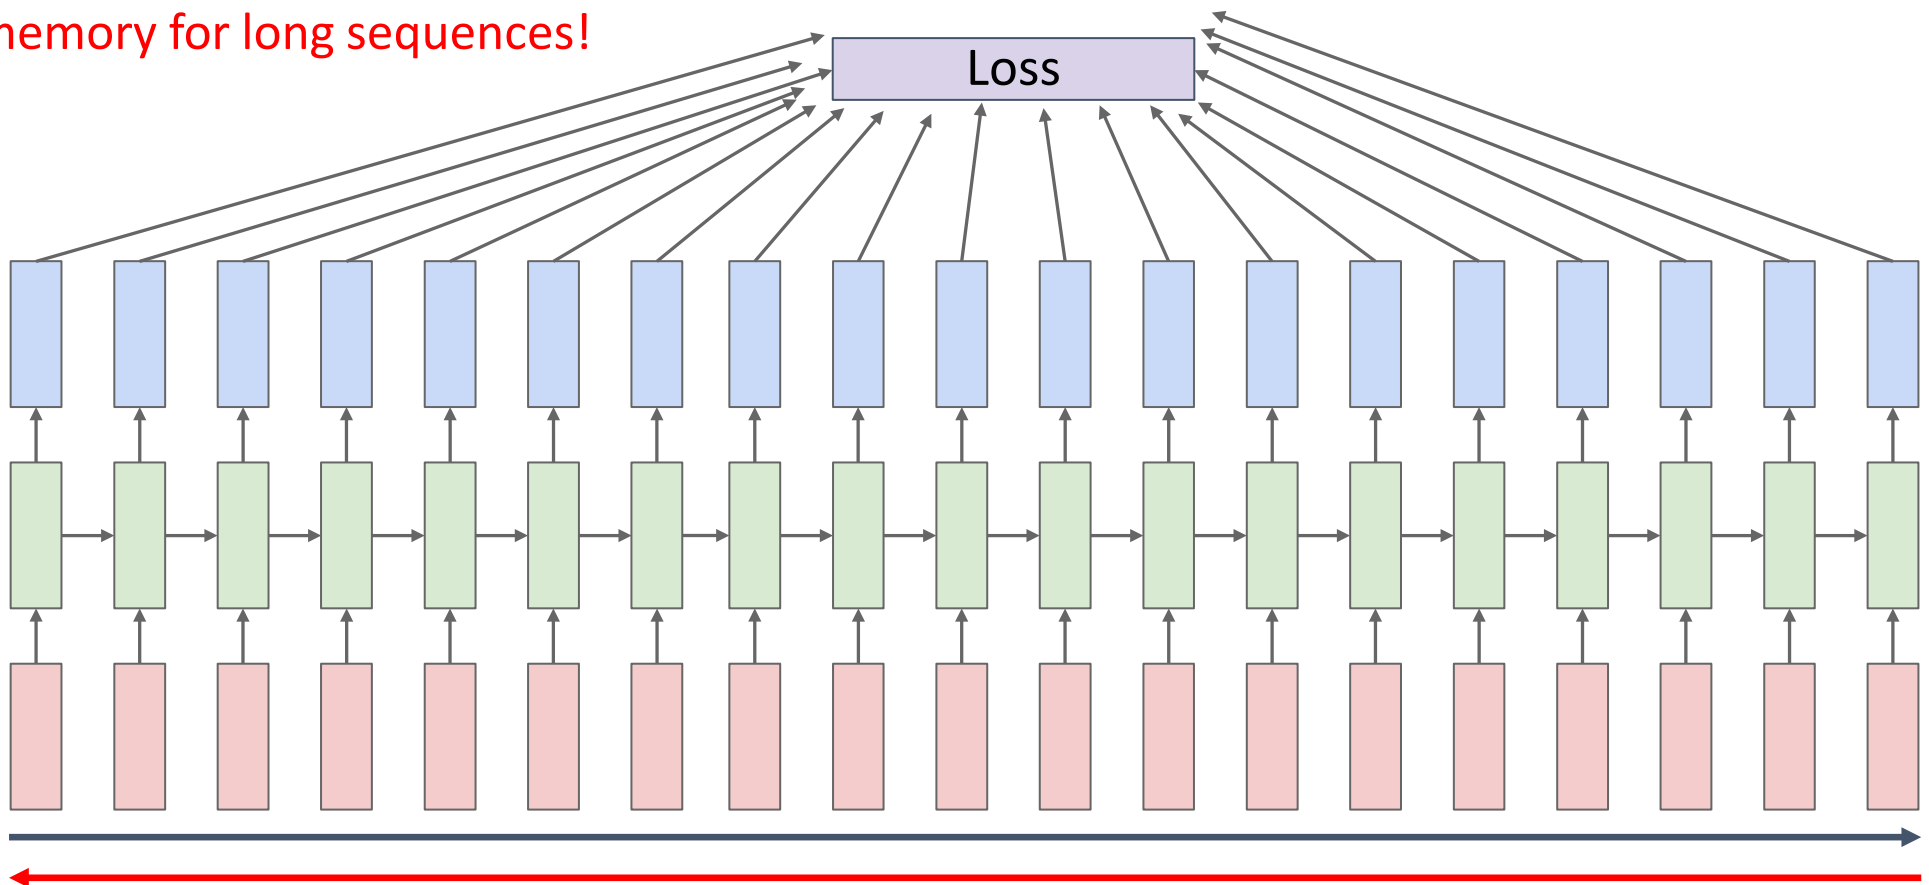
$W_{ih}$




# Backpropagation Through Time

Forward through entire sequence to compute loss, then backward through entire sequence to compute gradient

Problem: Takes a lot of memory for long sequences!



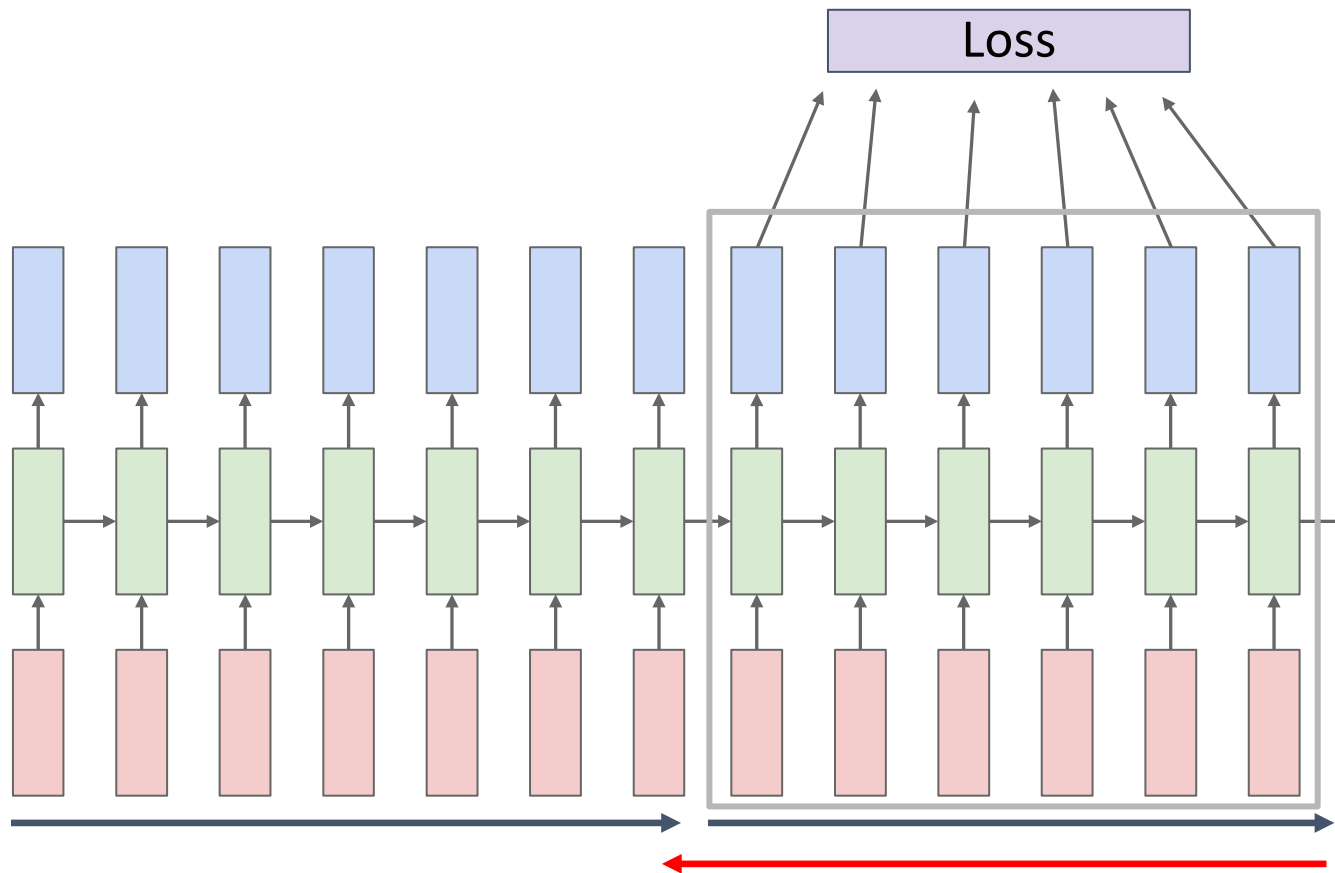
# Truncated Backpropagation Through Time



Run forward and backward through chunks of the sequence instead of whole sequence




# Truncated Backpropagation Through Time



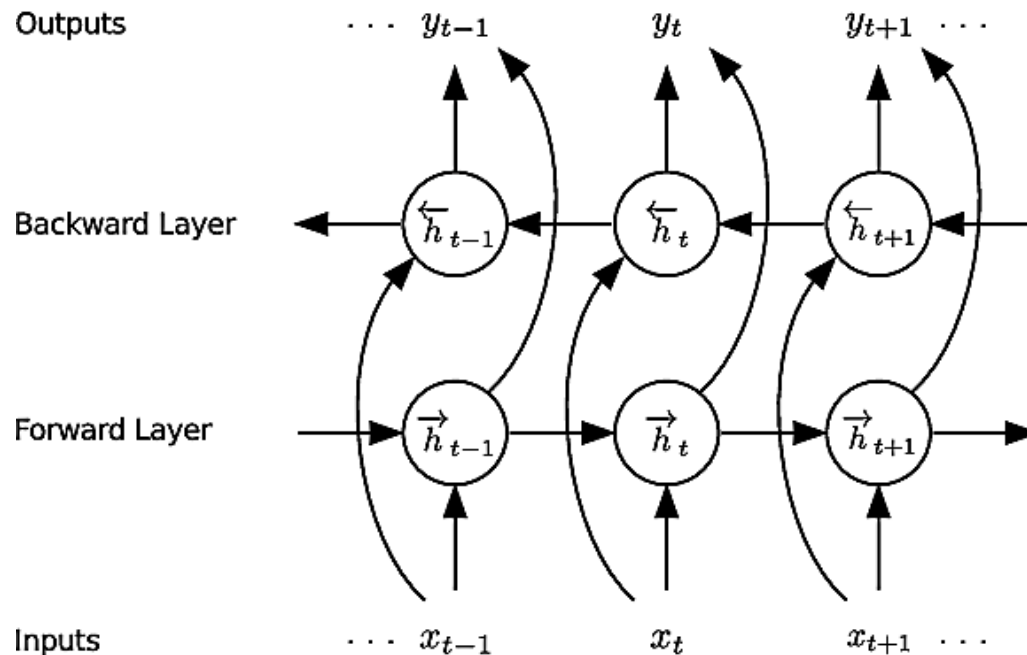
Carry hidden states forward in time forever, but only backpropagate for some smaller number of steps

# Truncated Backpropagation Through Time



# Bidirectional RNNs

- **Bidirectional RNNs** incorporate both forward and backward passes through sequential data
  - The output may not only depend on the previous elements in the sequence, but also on future elements in the sequence
  - It resembles two RNNs stacked on top of each other



$$\vec{h}_t = \sigma(\vec{W}^{(hh)}\vec{h}_{t-1} + \vec{W}^{(hx)}x_t)$$

$$\overleftarrow{h}_t = \sigma(\overleftarrow{W}^{(hh)}\overleftarrow{h}_{t+1} + \overleftarrow{W}^{(hx)}x_t)$$

$$y_t = f([\vec{h}_t; \overleftarrow{h}_t])$$

Outputs both past and future elements

# A recurrent neural language model

■ How does this compare to n-gram models?

## Improvements:

- Model size:  $O(V)$ , not  $O(V^n)$
- Sparsity (lack thereof)
- Sharing of representations across words
- Models long context

## Remaining challenges:

# A recurrent neural language model

■ How does this compare to n-gram models?

## Improvements:



- Model size:  $O(V)$ , not  $O(V^n)$
- Sparsity (lack thereof)
- Sharing of representations across words
- Models long context

## Remaining challenges:

- Softmax over large vocabulary
- High variance / overfitting
- Exploding and vanishing gradients

# But what about that huge softmax over $V$ ?

- Same problem as word embeddings: don't want to score wrt entire vocab!
- Solutions:
  - Hierarchical softmax
  - Noise-contrastive estimation (NCE)
  - Adaptive softmax [[Grave et al. 2017](#)]



# Problems learning RNNs

- In theory, should be able to propagate information over arbitrarily long contexts.
- In practice, RNNs suffer from **vanishing gradients** that decay to 0, or **exploding gradients** that increase towards infinity.

# Problems learning RNNs

- In theory, should be able to propagate information over arbitrarily long contexts.
- In practice, RNNs suffer from **vanishing gradients** that decay to 0, or **exploding gradients** that increase towards infinity.
  - Exploding gradients mostly resolved by **gradient clipping**: thresholding gradient values, or rescaling them. Threshold/scale is a hyperparameter.

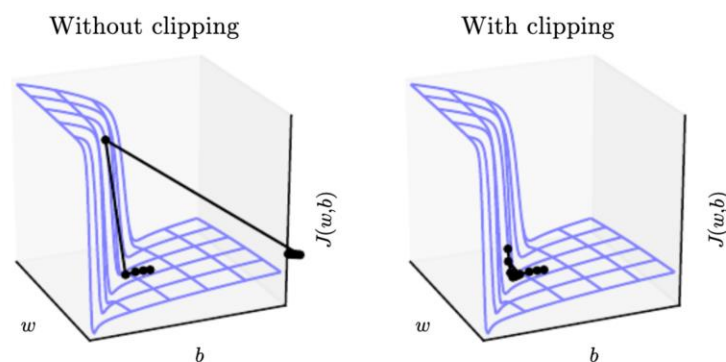


Figure: Goodfellow, Bengio and Courville. Deep Learning. MIT Press, 2016.



# Problems learning RNNs

- In theory, should be able to propagate information over arbitrarily long contexts.
- In practice, RNNs suffer from **vanishing gradients** that decay to 0, or **exploding gradients** that increase towards infinity.
  - Exploding gradients mostly resolved by **gradient clipping**: thresholding gradient values, or rescaling them. Threshold/scale is a hyperparameter.
- Vanishing gradients mostly resolved by adding **gating** to the RNN composition function.
  - Sigmoid activation function in RNN leads to this problem.
  - Relu, in theory, avoids this problem but not in practice.

